

PROCESS MANAGEMENT

Chapter 3: Process Concept

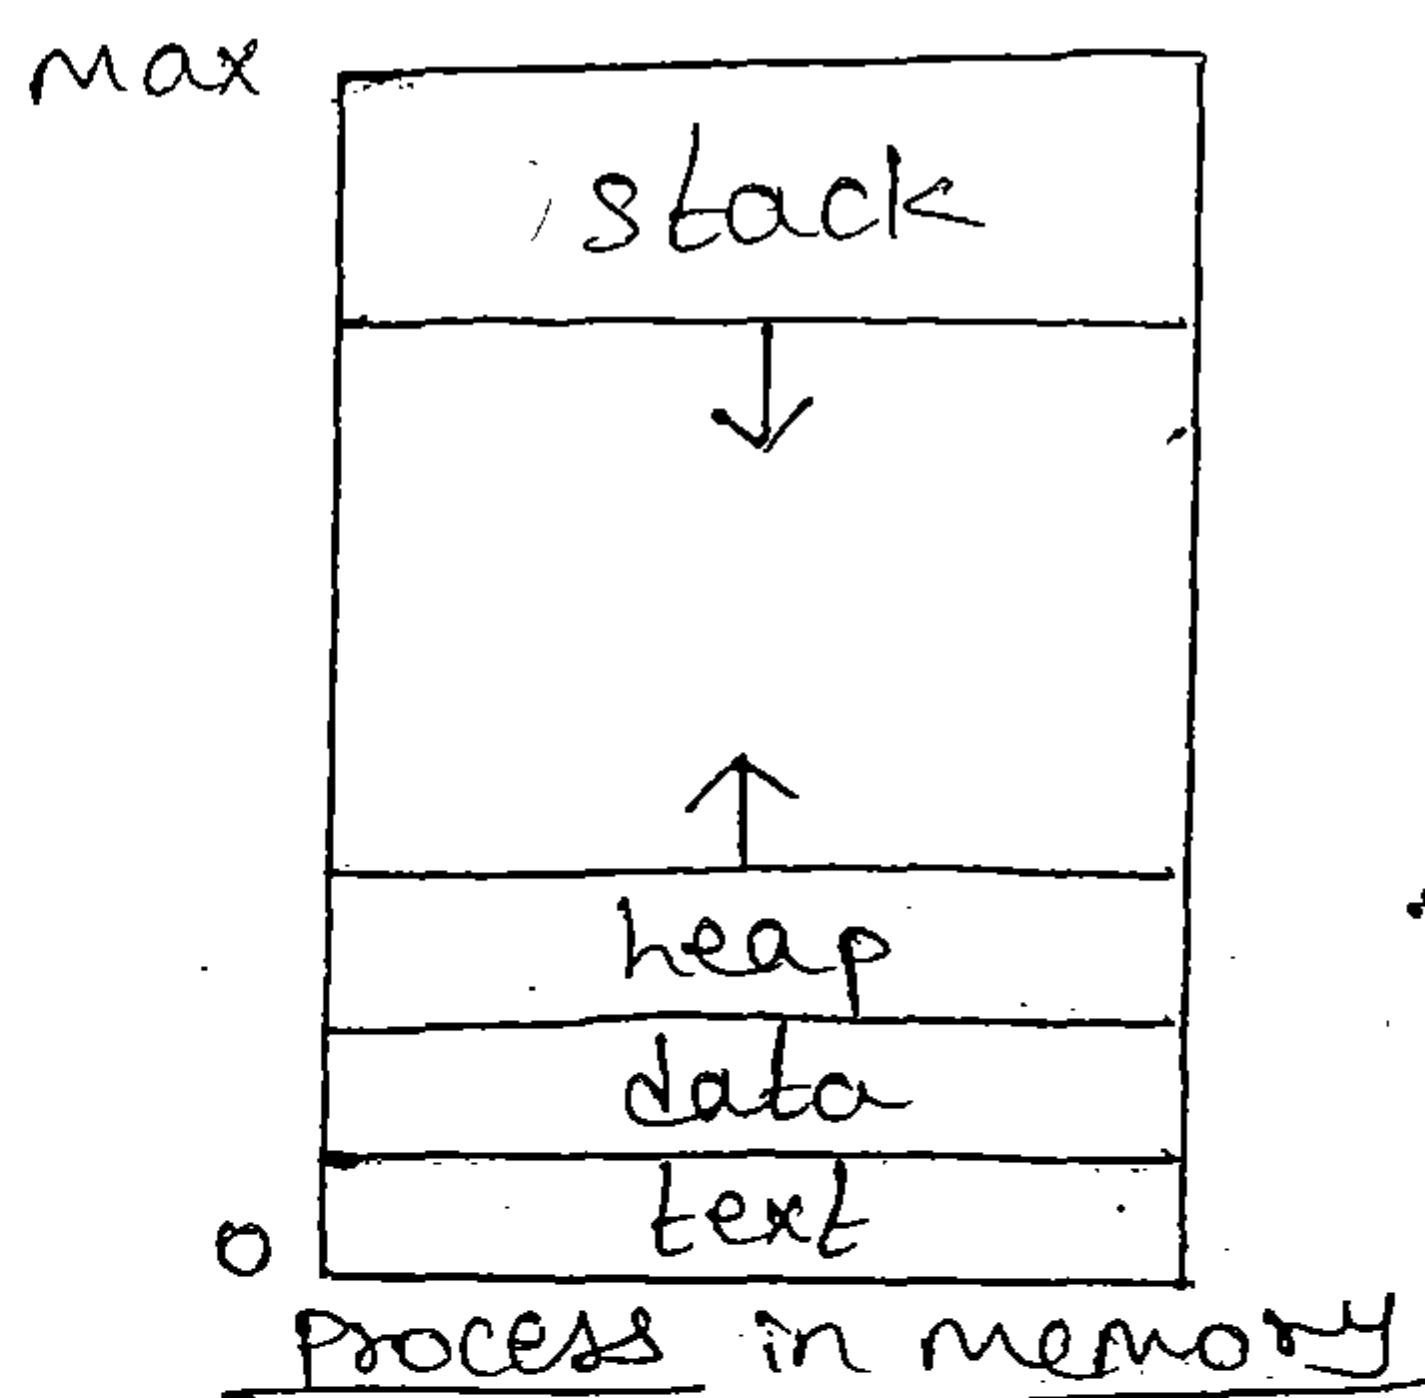
- Objectives:
- * To introduce the notion of a process - a program in execution, which forms the basis of all computation.
 - * To describe the various features of processes, including scheduling, creation and termination, and communication.
 - * To describe communication in client-server systems.

⑩ Process Concept

- An OS executes a variety of programs:
 - * A batch system executes jobs.
 - * A time-shared system has user programs (or) tasks.
- Process - A program in execution; process execution must progress in sequential fashion.
- A process includes:
 - * program counter
 - * stack
 - * data section

The Process

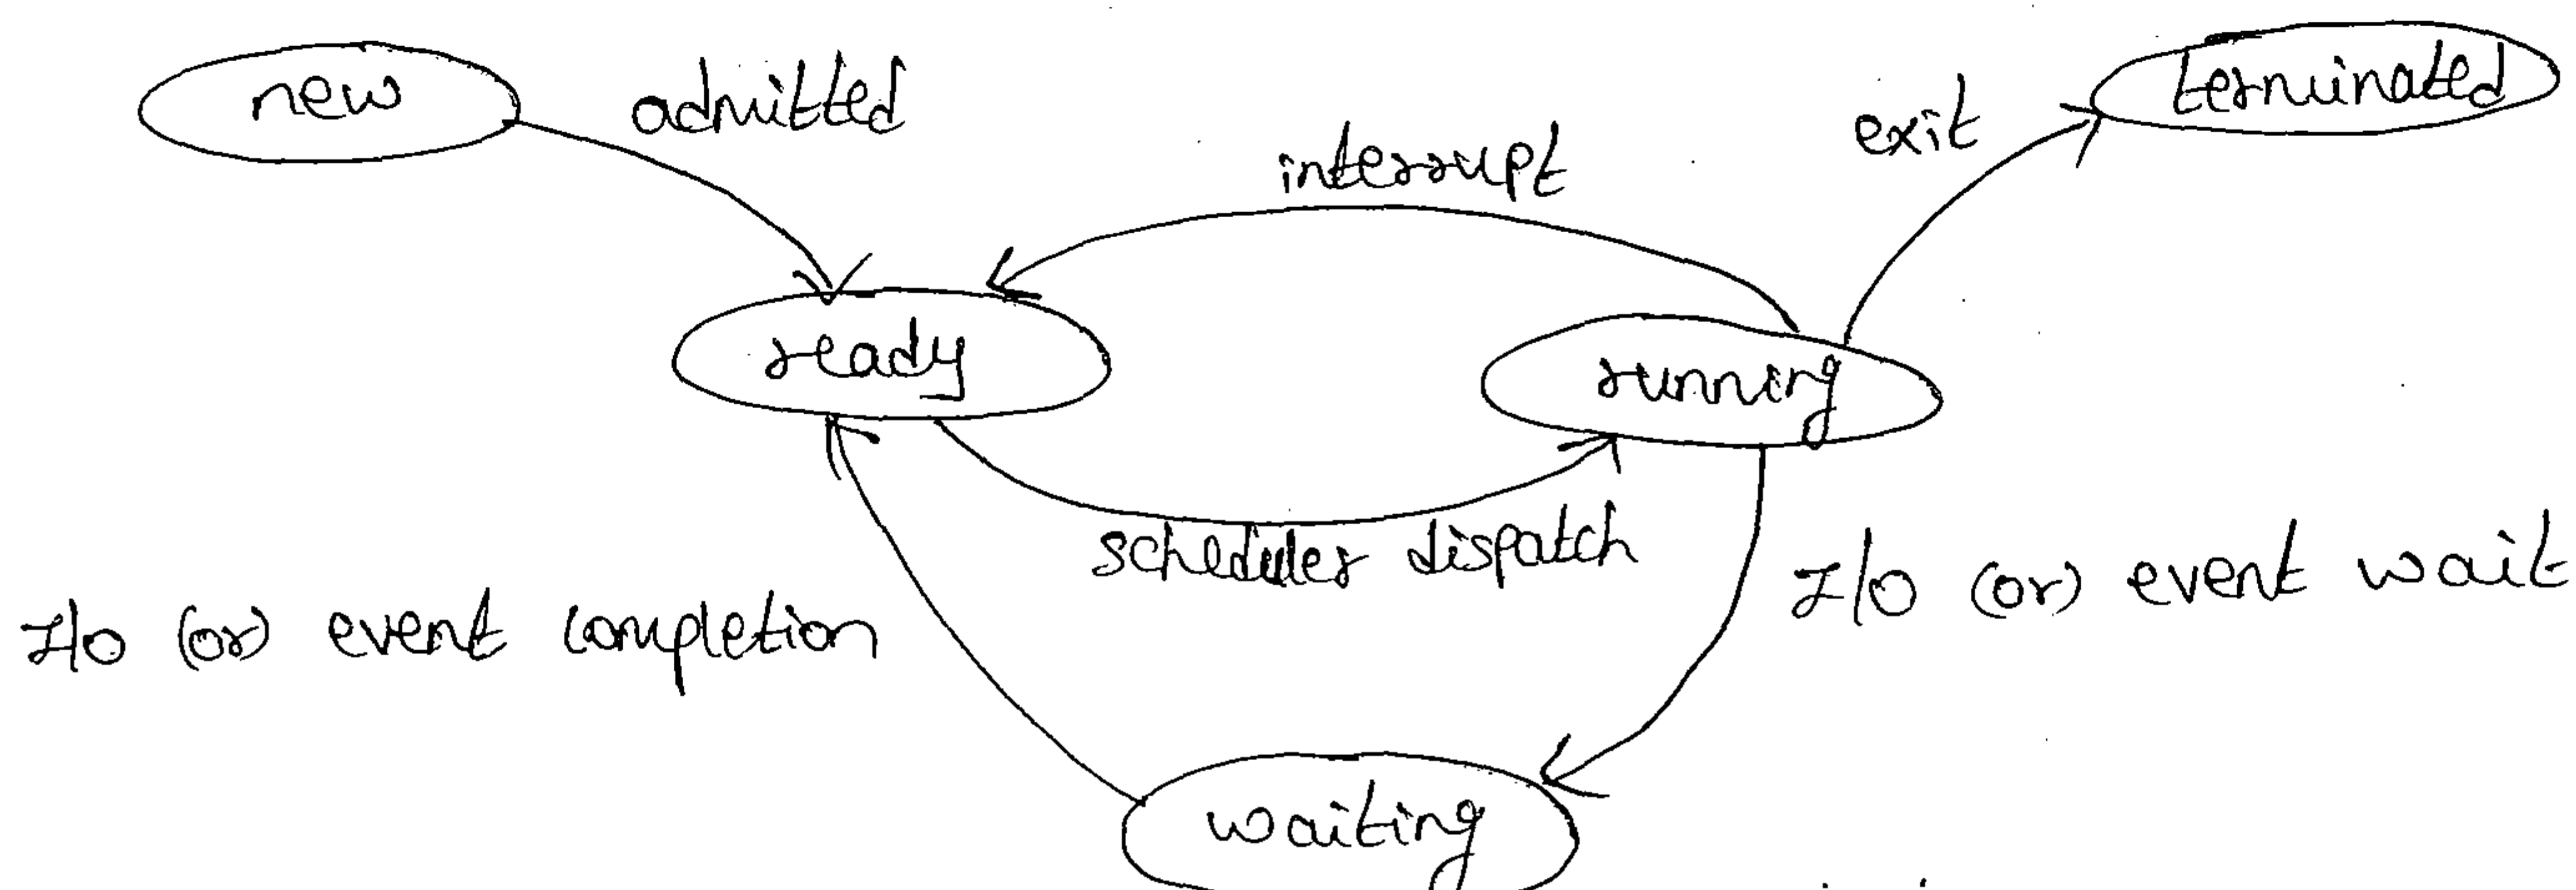
- Multiple parts:
 - * The program code, also called text section.
 - * Current activity including program counter, processor registers.
 - * Stack containing temporary data
 - function parameters, return addresses, local variables
 - * Data section containing global variables.



- > Program is passive entity, process is active entity.
- > Program becomes process when executable file loaded into memory.
- > Execution of program started via GUI mouse clicks, command line entry of its name etc.
- > One program can be several processes.
 - * Consider multiple users executing the same program.

Process State

- > As a process executes, it changes state.
 - 1) New - The process is being created.
 - 2) Running - Instructions are being executed.
 - 3) Waiting - The process is waiting for some event to occur.
 - 4) Ready - The process is waiting to be assigned to a processor.
 - 5) Terminated - The process has finished execution.



44 process control block


- Each process is represented in the OS by a process control block (PCB).
- Also called a task control block.

Process state
Process number
Program counter
registers
memory limits
list of open files
...

Process control block (PCB)

- PCB contains many pieces of information associated with a specific process:

- Process state
- Program Counter
- CPU registers
- CPU-scheduling information
- Memory-management information
- Accounting information
- I/O status information



① Process Scheduling

42

- Maximize CPU use, quickly switch processes onto CPU for time sharing.
- Process scheduler selects among available processes for next execution on CPU.

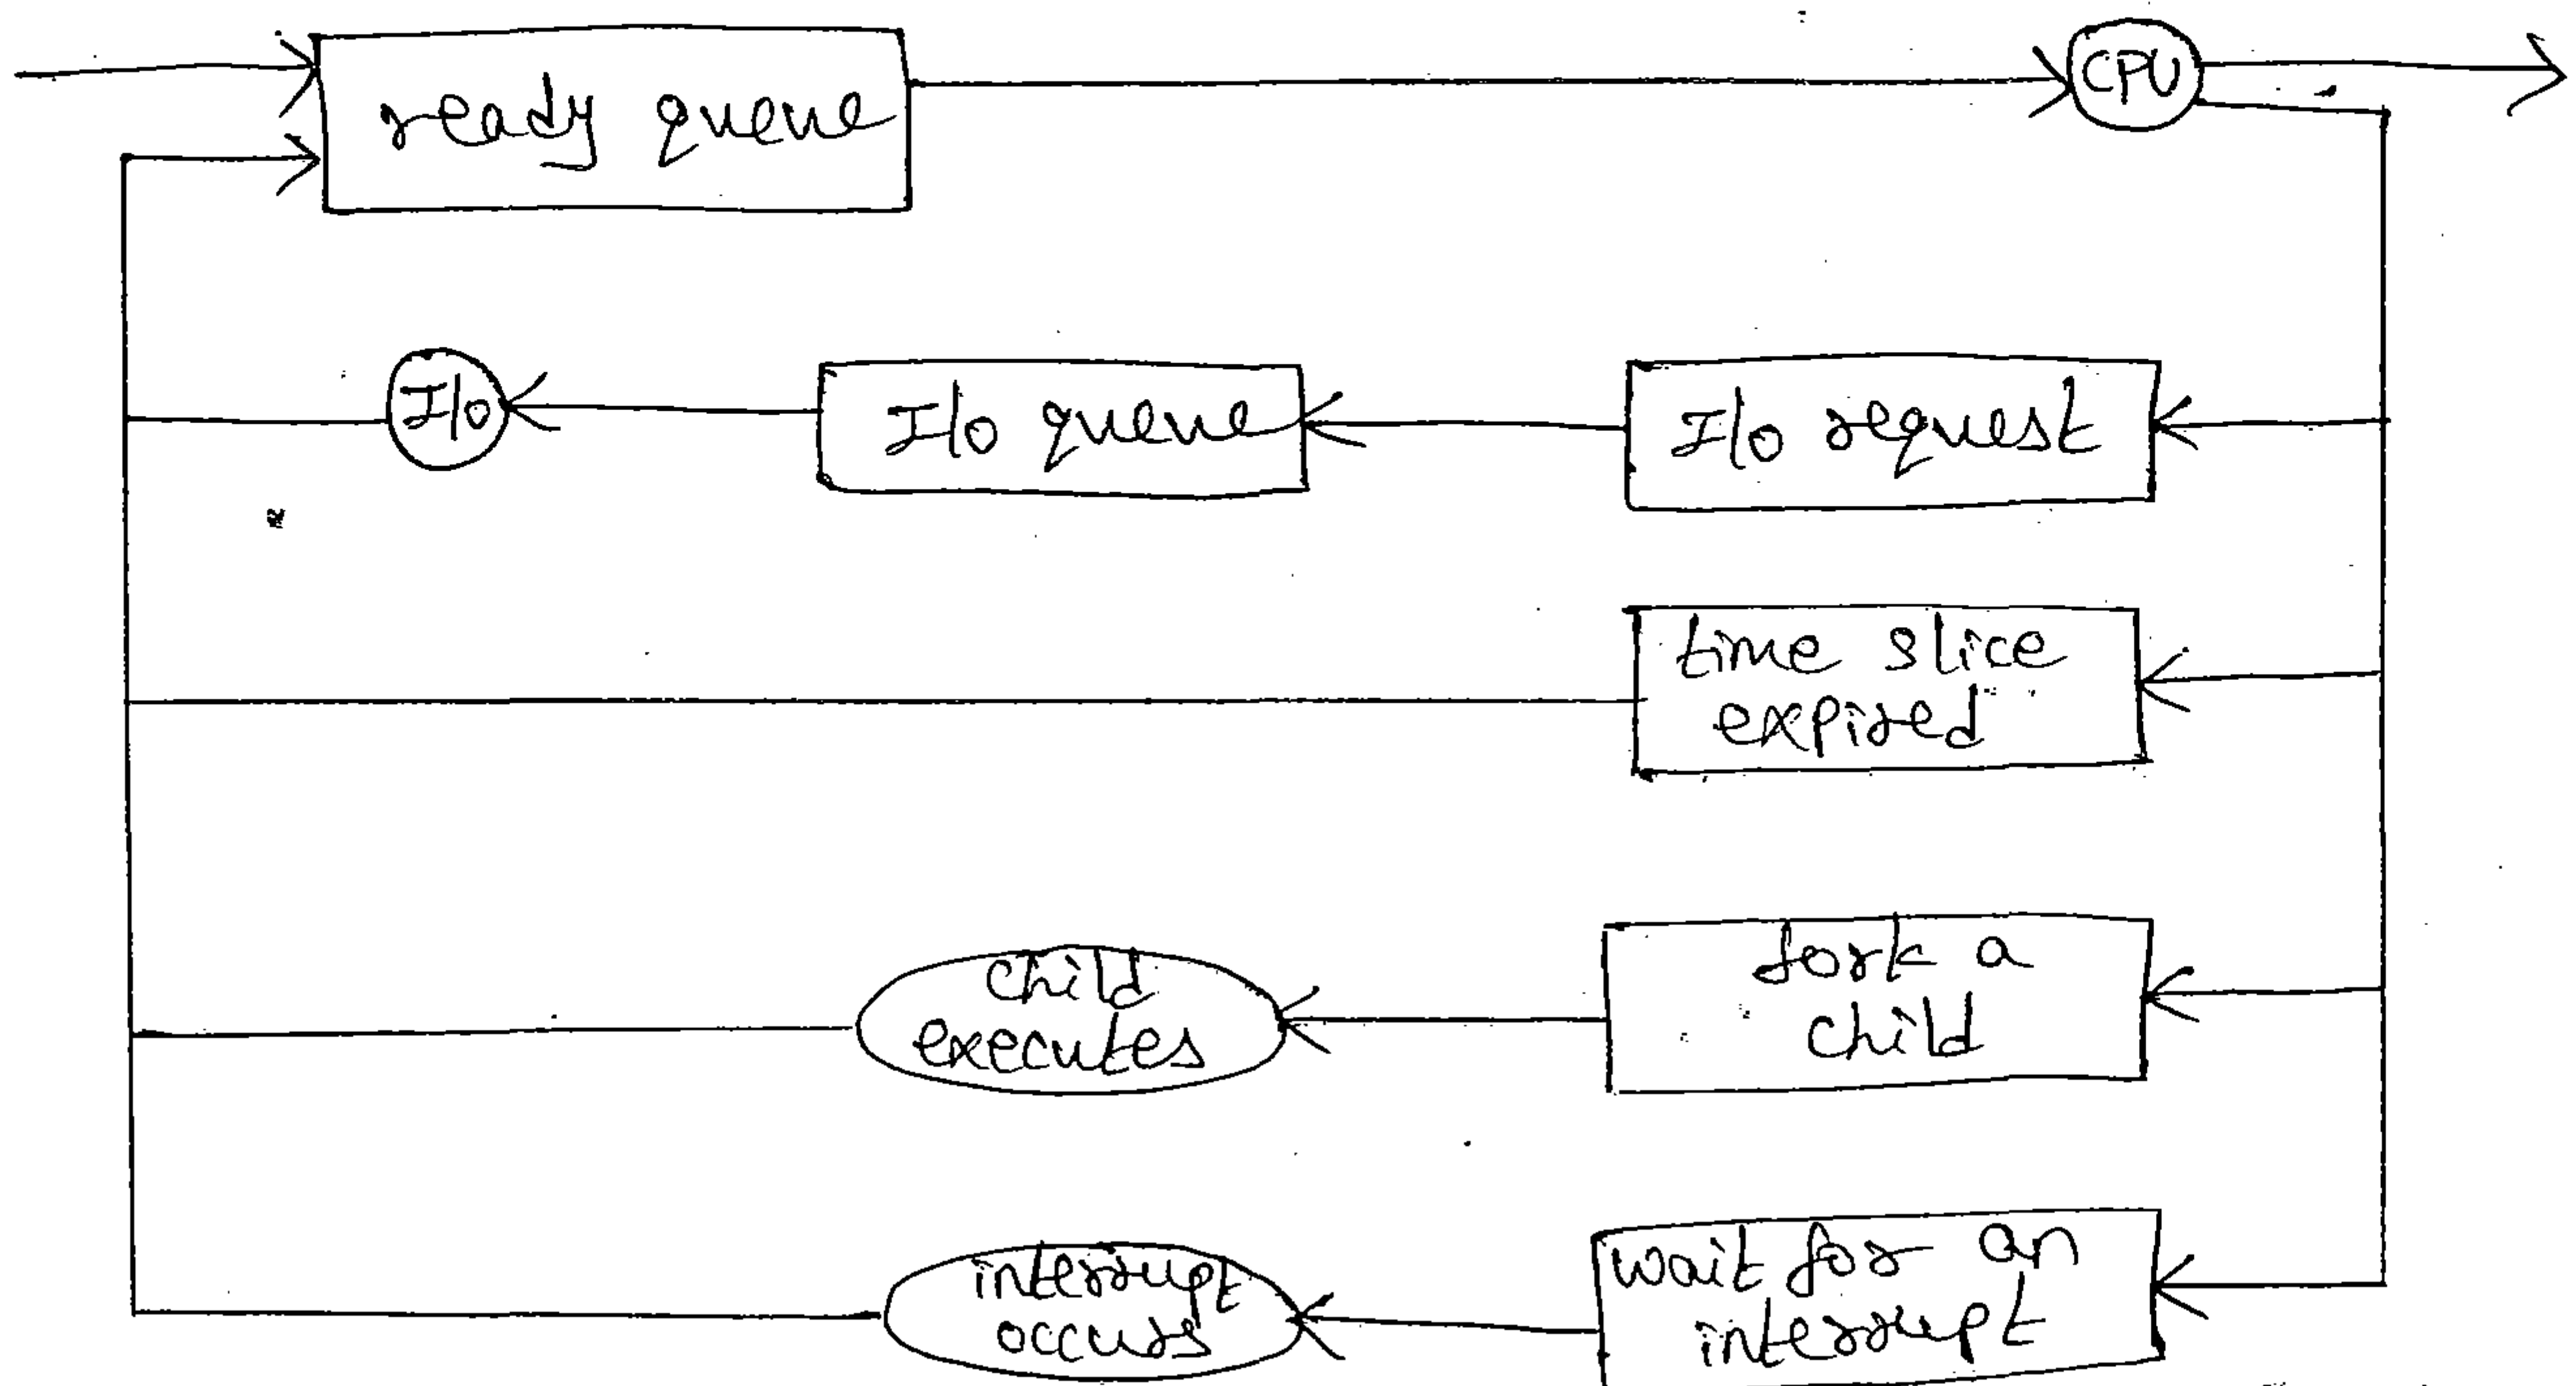
Scheduling Queues

- * Job queue - set of all processes in the system.
- * Ready queue - set of all processes residing in main memory, ready & waiting to execute. [Refer Figure 3.6]
- * Device queues - set of processes waiting for an I/O device.

→ Processes migrate among the various queues.

→ A common representation of process scheduling is a queuing diagram.

- * Each rectangular box represents a queue.
- * Two types of queues are present: ready queue, set of device queues.
- * Circles represent the resources that serve the queues.
- * Arrows indicate the flow of processes in the system.

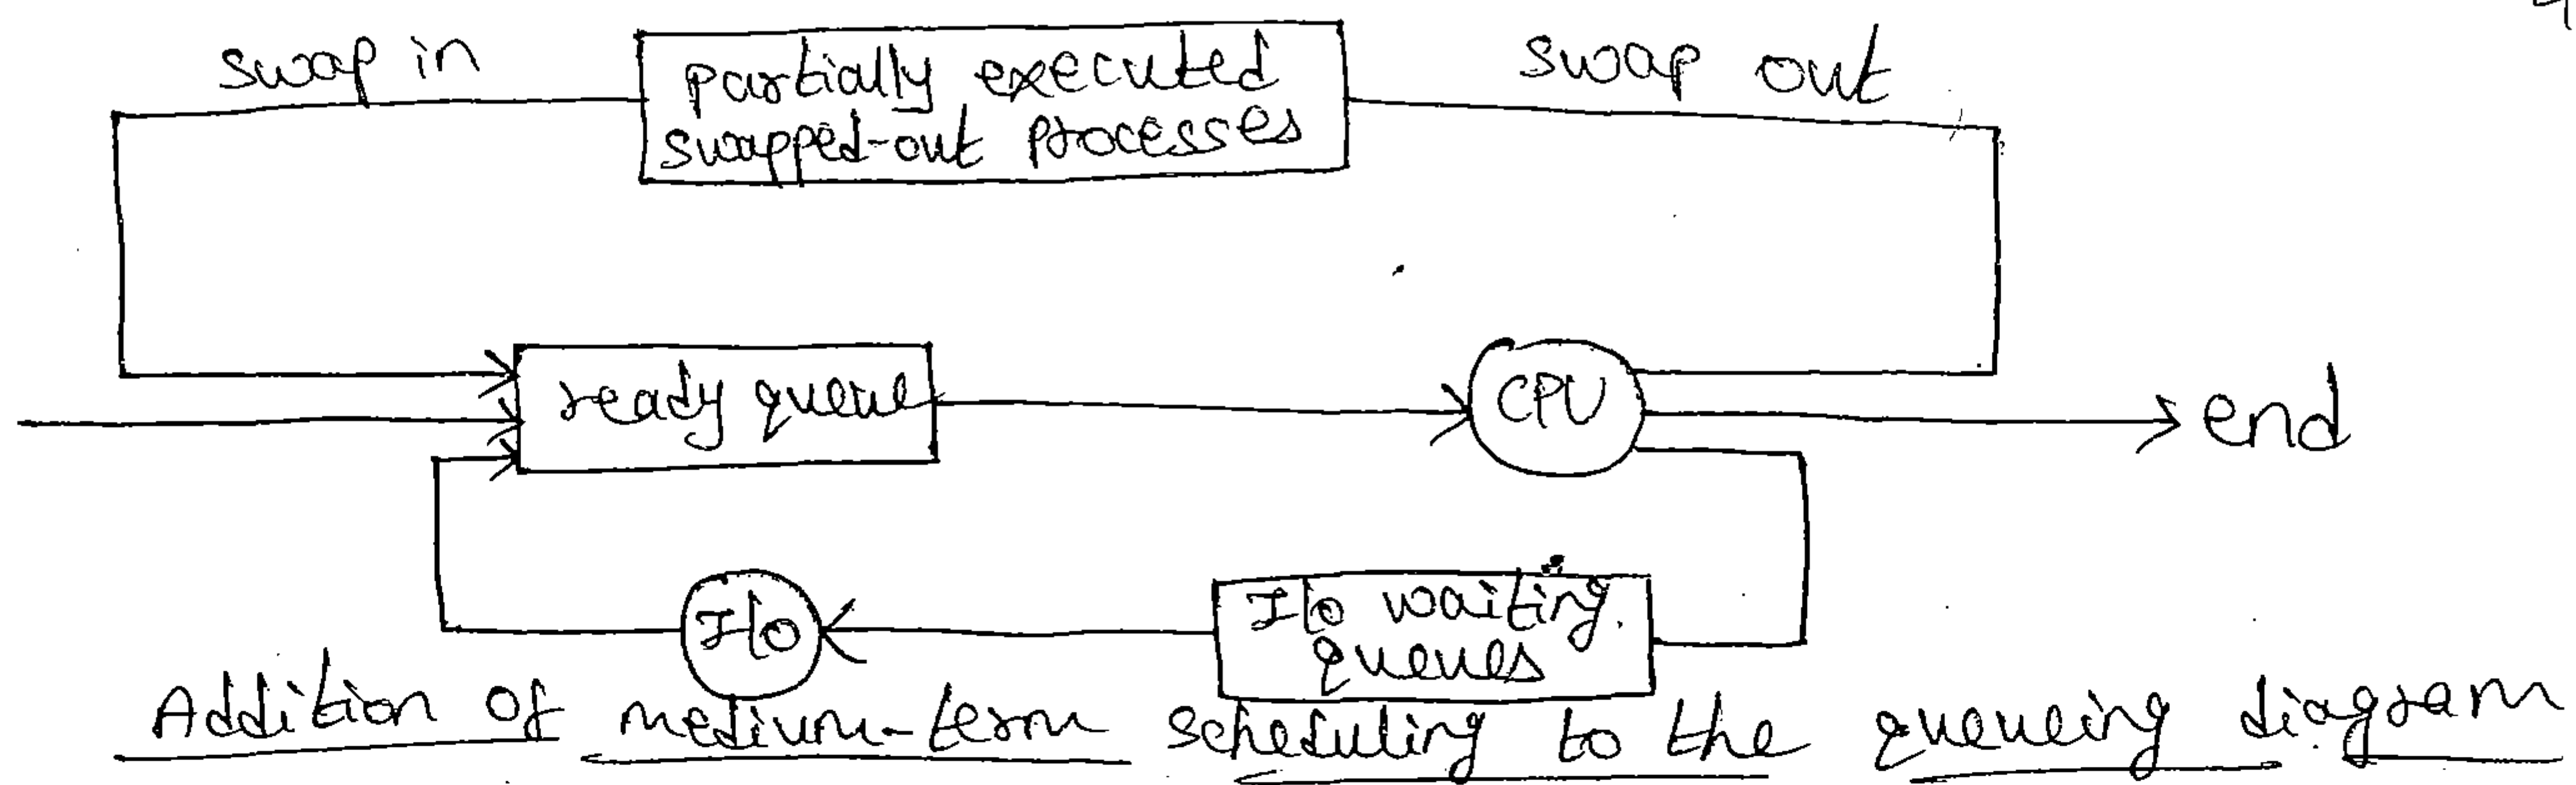


Queuing-Diagram representation of process scheduling

- A new process is initially put in a ready queue.
- It waits there until it is selected for execution (or) dispatched.
- Once the process is allocated the CPU & is executing, one of the several events could occur:
 - * The process could issue an I/O request and then be placed in an I/O queue.
 - * The process could create a new subprocess and wait for the subprocess's termination.
 - * The process could be removed forcibly from the CPU, as a result of an interrupt, & be put back in the ready queue.

Schedulers

- Long-term scheduler (or job scheduler) - selects which processes should be brought into the ready queue.
- Short-term scheduler (or CPU scheduler) - selects which process should be executed next and allocates CPU.
- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- The long-term scheduler controls the degree of multiprogramming (the number of processes in memory).
- Processes can be described as either:
 - * I/O-bound process - spends more time doing I/O than computations, many short CPU bursts.
 - * CPU-bound process - spends more time doing computations; few very long CPU bursts.
- Time-sharing systems, may introduce an additional intermediate level of scheduling - medium-term scheduler.
 - * Swapping.



Context Switch

- when CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a context switch.
- Context of a process represented in the PCB.
- Context-switch time is overhead; the system does no useful work while switching
 - * The more complex the OS & the PCB → longer the context switch.
- > Time dependent on hardware support
 - * Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once.

⑩ Operations on Processes

Process Creation

- The creating process is called a Parent process.
- Parent process create children processes, which in turn create other processes, forming a tree of processes.
- Generally, process identified and managed via a process identifier (Pid).
- Resource sharing:
 - * parent and children share all resources
 - * children share subset of parent's resources.

→ Execution: When a process creates a new process:- 4-

- (i) The parent continues to execute concurrently with its children.
- (ii) The parent waits until some (or) all of its children have terminated.

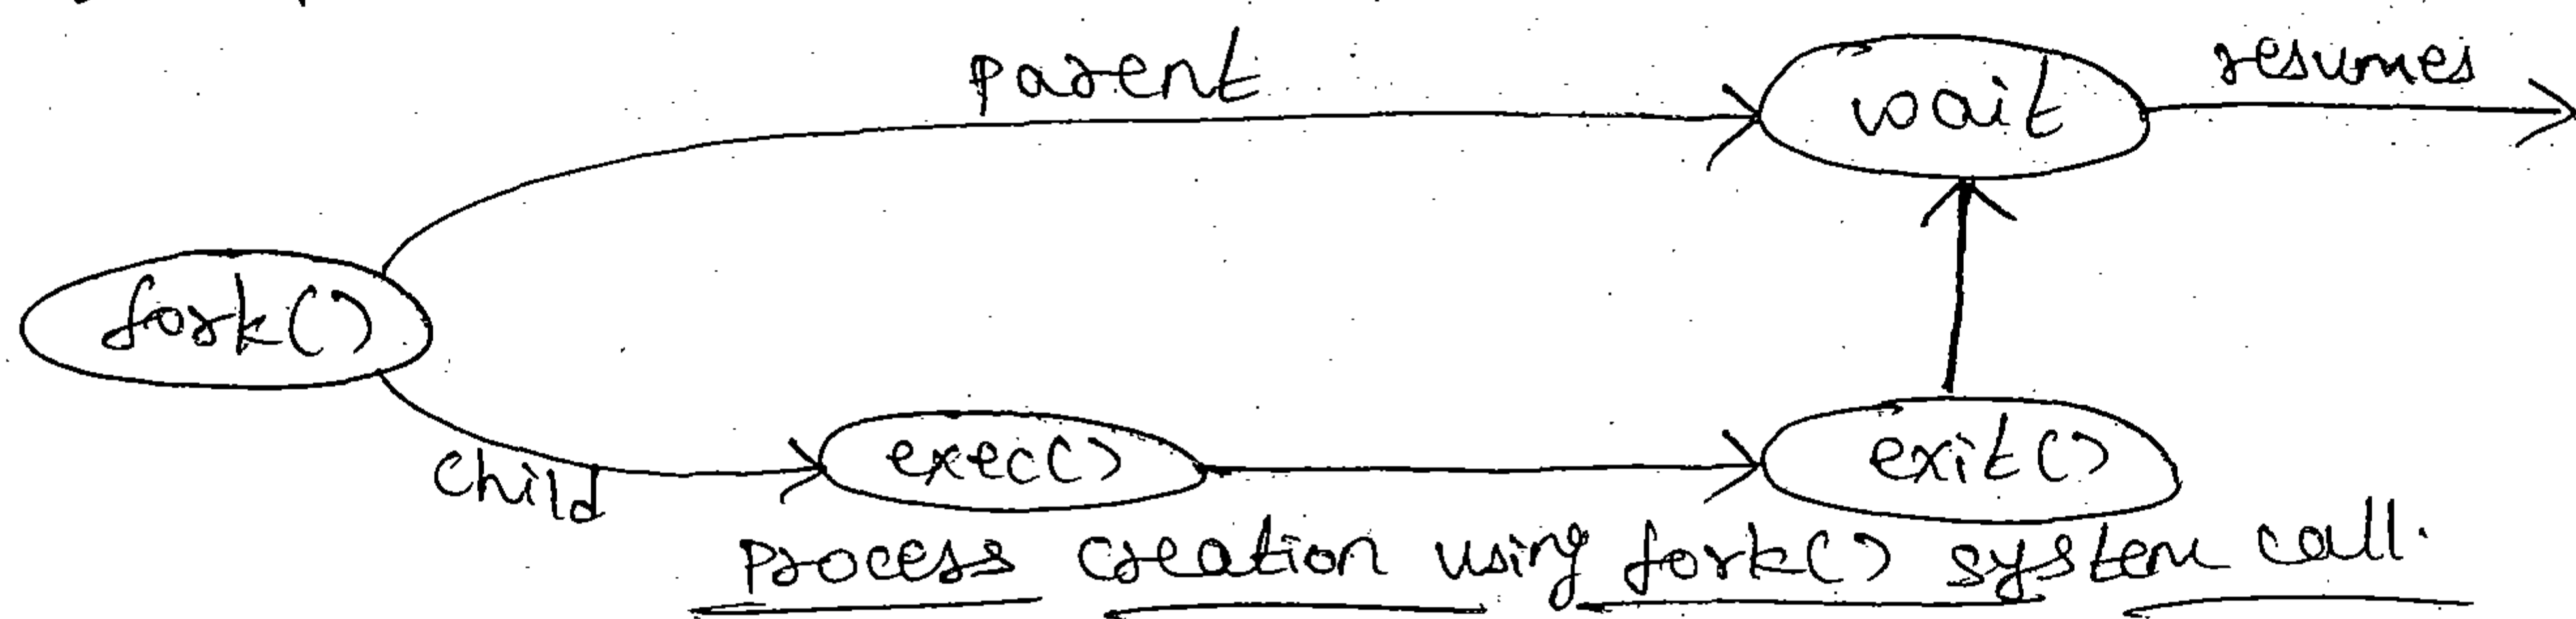
→ Also two possibilities in terms of address space of the new process:

- (i) The child process is a duplicate of the parent process.
- (ii) The child process has a new program loaded into it.

→ UNIX examples:

* `fork()` system call creates new process.

* `exec()` system call used after a `fork()` to replace the process' memory space with a new program.



Process Termination

→ Process executes last statement and asks the OS to delete it by using `exit()` system call.

→ Output data from child to parent (via `wait()` system call).

→ Process' resources are deallocated by OS.

→ Parent may terminate execution of children processes (via `abort()` system call).

* Child has exceeded allocated resources

* Task assigned to child is no longer required.

* If parent is exiting

→ Some OS do not allow child to continue if its parent terminates.



- All children terminated: Cascading Termination.

⑦ Interprocess Communication

- A process is independent if it cannot affect (or) be affected by the other processes executing in the system.
- Cooperating process can affect (or) be affected by other processes, including sharing data.
- Reasons for providing an environment that allows process cooperation:

- * Information sharing
- * Computation speedup
- * Modularity
- * Convenience

- Cooperating processes need interprocess communication (IPC).
- Two models of IPC:
 - (i) Shared memory
 - (ii) Message passing



Communications models. (a) message passing. (b) shared memory.

Shared-memory Systems

- A region of memory that is shared by cooperating processes is established.
- Processes can then exchange information by reading & writing

→ Producer-Consumer Problem

* Paradigm for cooperating processes.

* Producer process produces information that is consumed by a consumer process.

* Example: A compiler may produce assembly code, which is consumed by an assembler. The assembler, in turn, may produce object modules, which are consumed by the loader.

* One solution to the producer-consumer problem uses shared memory.

* Two types of buffers:

» Unbounded-buffer places no practical limit on the size of the buffer.

» Bounded-buffer assumes that there is a fixed buffer size.

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    ...
    item;
```

```
    item buffer[BUFFER_SIZE];
```

```
    int in = 0;
```

```
    int out = 0;
```

* Solution is correct, but can only use $BUFFER_SIZE - 1$ elements.

→ The code for the producer and consumer processes:

```
    item nextProduced;
```

```
    while (true) {
```

```
        /* Produce an item in nextProduced */
```

```
        while ((in + 1) % BUFFER_SIZE == out)
```

```
            ; /* do nothing */
```

```

    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}

```

The Producer process

```

item nextConsumed;
while (true) {
    while (in == out)
        ; // do nothing
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in nextConsumed */
}

```

The Consumer process

Message-Passing Systems

- > Mechanism for processes to communicate and to synchronize their actions in distributed environment.
- > message system - processes communicate with each other without resorting to shared variables/same address space.
- > A message-passing facility provides at least two operations:
 - * send(message) - message size fixed (or) variable.
 - * receive(message) Ex:- chat program used on the WWW.
- > If Processes P and Q want to communicate, they need to:
 - * Establish a communication link between them.
 - * Exchange messages via send/receive.
- > Implementation of communication link:
 - * physical (Ex:- shared memory, hardware bus)
 - * logical (Ex:- logical partition)

→ Methods for logically implementing a link and the send() receive() operations:

- Direct (or) indirect communication
- Synchronous (or) asynchronous communication.
- Automatic (or) explicit buffering.

→ Naming

* Under direct communication, each process that wants to communicate must explicitly name the recipient (or) sender of the communication.

→ send(P, message) - Send a message to process P.

→ receive(Q, message) - Receive a message from process Q.

* Properties of communication link:

→ Links are established automatically

→ A link is associated with exactly one pair of communicating processes.

→ Between each pair there exists exactly one link.

→ The link may be unidirectional, but is usually bi-directional.

* With indirect communication, messages are directed and received from mailboxes (also referred to as ports)

→ Each mailbox has a unique id.

→ Processes can communicate only if they share a mailbox.

→ send(A, message) - Send a message to mailbox A.

→ receive(A, message) - Receive a message from mailbox A.

* Properties of communication link:

→ Link established only if processes share a common mailbox.

→ A link may be associated with many processes.

50
⇒ Each pair of processes may share several communication links.

⇒ Link may be unidirectional (or) bi-directional.

* Operations:

⇒ Create a new mailbox.

⇒ Send and receive messages through the mailbox.

⇒ Delete a mailbox.

* Mailbox sharing:

⇒ P_1, P_2 & P_3 share mailbox A

⇒ P_1 sends; P_2 and P_3 receive

⇒ Which process will receive the message sent by P_1 ?

* Solutions:

⇒ Allow a link to be associated with at most two processes.

⇒ Allow only one process at a time to execute a receive operation.

⇒ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

→ Synchronization

* Message passing may be either blocking (or) non-blocking.

* Blocking is considered as synchronous.

⇒ Blocking send has the sender block until the message is received.

⇒ Blocking receive has the receiver block until a message is available.

* Non-blocking is considered as asynchronous

⇒ Non-blocking send has the sender send the message and continue.

→ Non-blocking receive has the receiver receive a valid message (or) null.

→ Buffering

* Queue of messages attached to the link; implemented in one of three ways:

(i) Zero capacity - queue has a maximum length of 0.
sender must wait for receiver.
(block)

(ii) Bounded capacity - queue has finite length 'n'.
sender must wait if link is full.
(block)

(iii) Unbounded capacity - queue's length is potentially infinite. The sender never blocks.

Chapter 4: Multithreaded Programming

Objectives: * To introduce the notion of a thread - a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems.

* To discuss the APIs for the Pthreads, Win32 and Java thread libraries.

* To examine issues related to multithreaded programming.

① Overview

→ A thread is a basic unit of CPU utilization.

→ A thread is a single sequence stream within in a process.

→ A traditional (or) heavyweight process has a single thread of control.

→ Multiple tasks with the application can be implemented by separate threads:

* Update display

* Fetch data

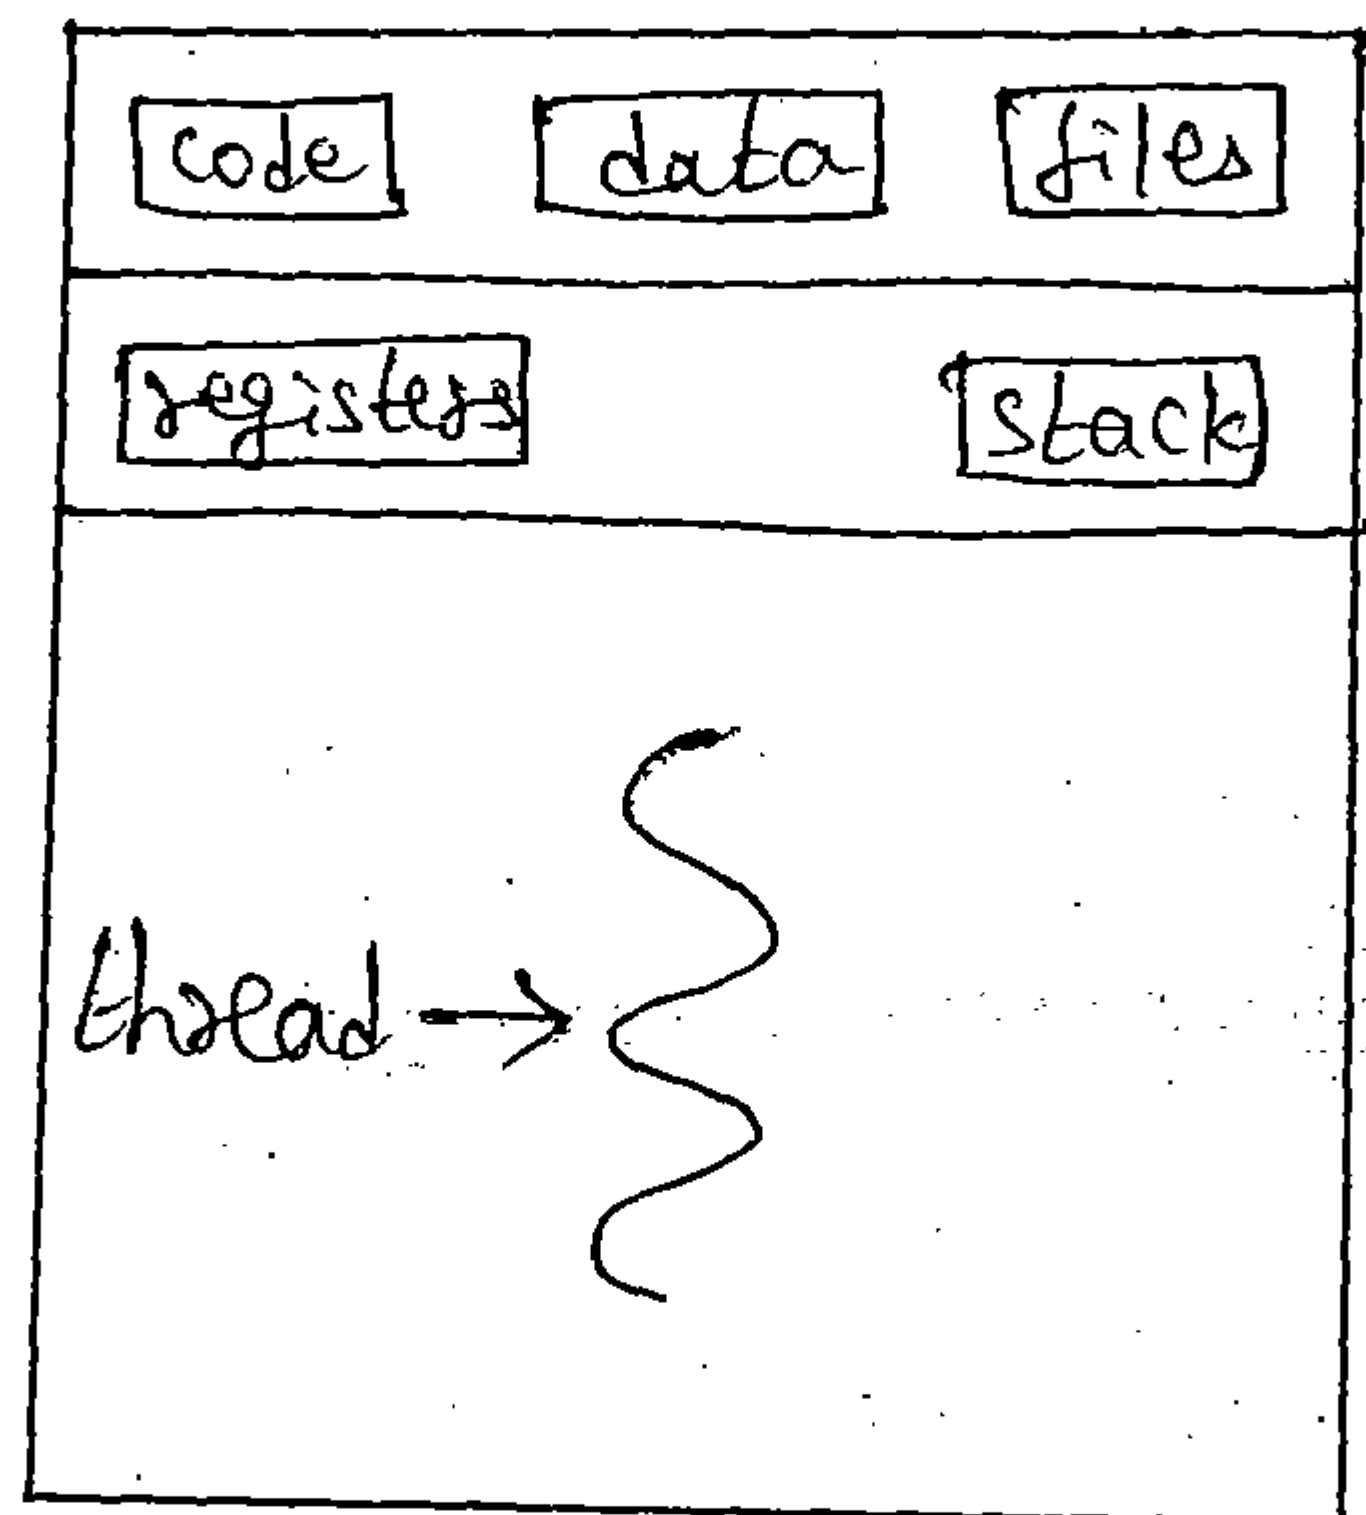
* Spell checking

... and so on.


52

→ Process creation is heavy-weight while thread creation is light-weight.


→ kernels are generally multithreaded.



Single-threaded process





Multi-threaded process



* Benefits

- * **Responsiveness** - Multithreading an interactive application may allow a program to continue running even if part of it is blocked (or) is performing a lengthy operation, thereby increasing responsiveness to the user.
- * **Resource sharing** - Processes may only share resources through techniques such as shared memory (or) message passing.
- * **Economy** - Allocating memory and resources for process creation is costly. More economical to create & context-switch threads.
- * **Scalability** - The benefits of multithreading can be greatly increased in a multiprocessor architecture, where threads can be running in parallel on different processors.

Multicore Programming



* Multicore systems putting pressure on programmers, challenges include:


- (i) Dividing activities
- (ii) Balance
- (iii) Data splitting
- (iv) Data dependency
- (v) Testing and debugging

① Multithreading Models

→ Support for threads may be provided either at the user level, for user threads; or by the kernel, for kernel threads.


→ kernel threads examples: Windows XP, Linux, Mac OS X, Solaris and Tru64 UNIX (formerly, Digital UNIX) - support kernel threads.

Many-to-one Model



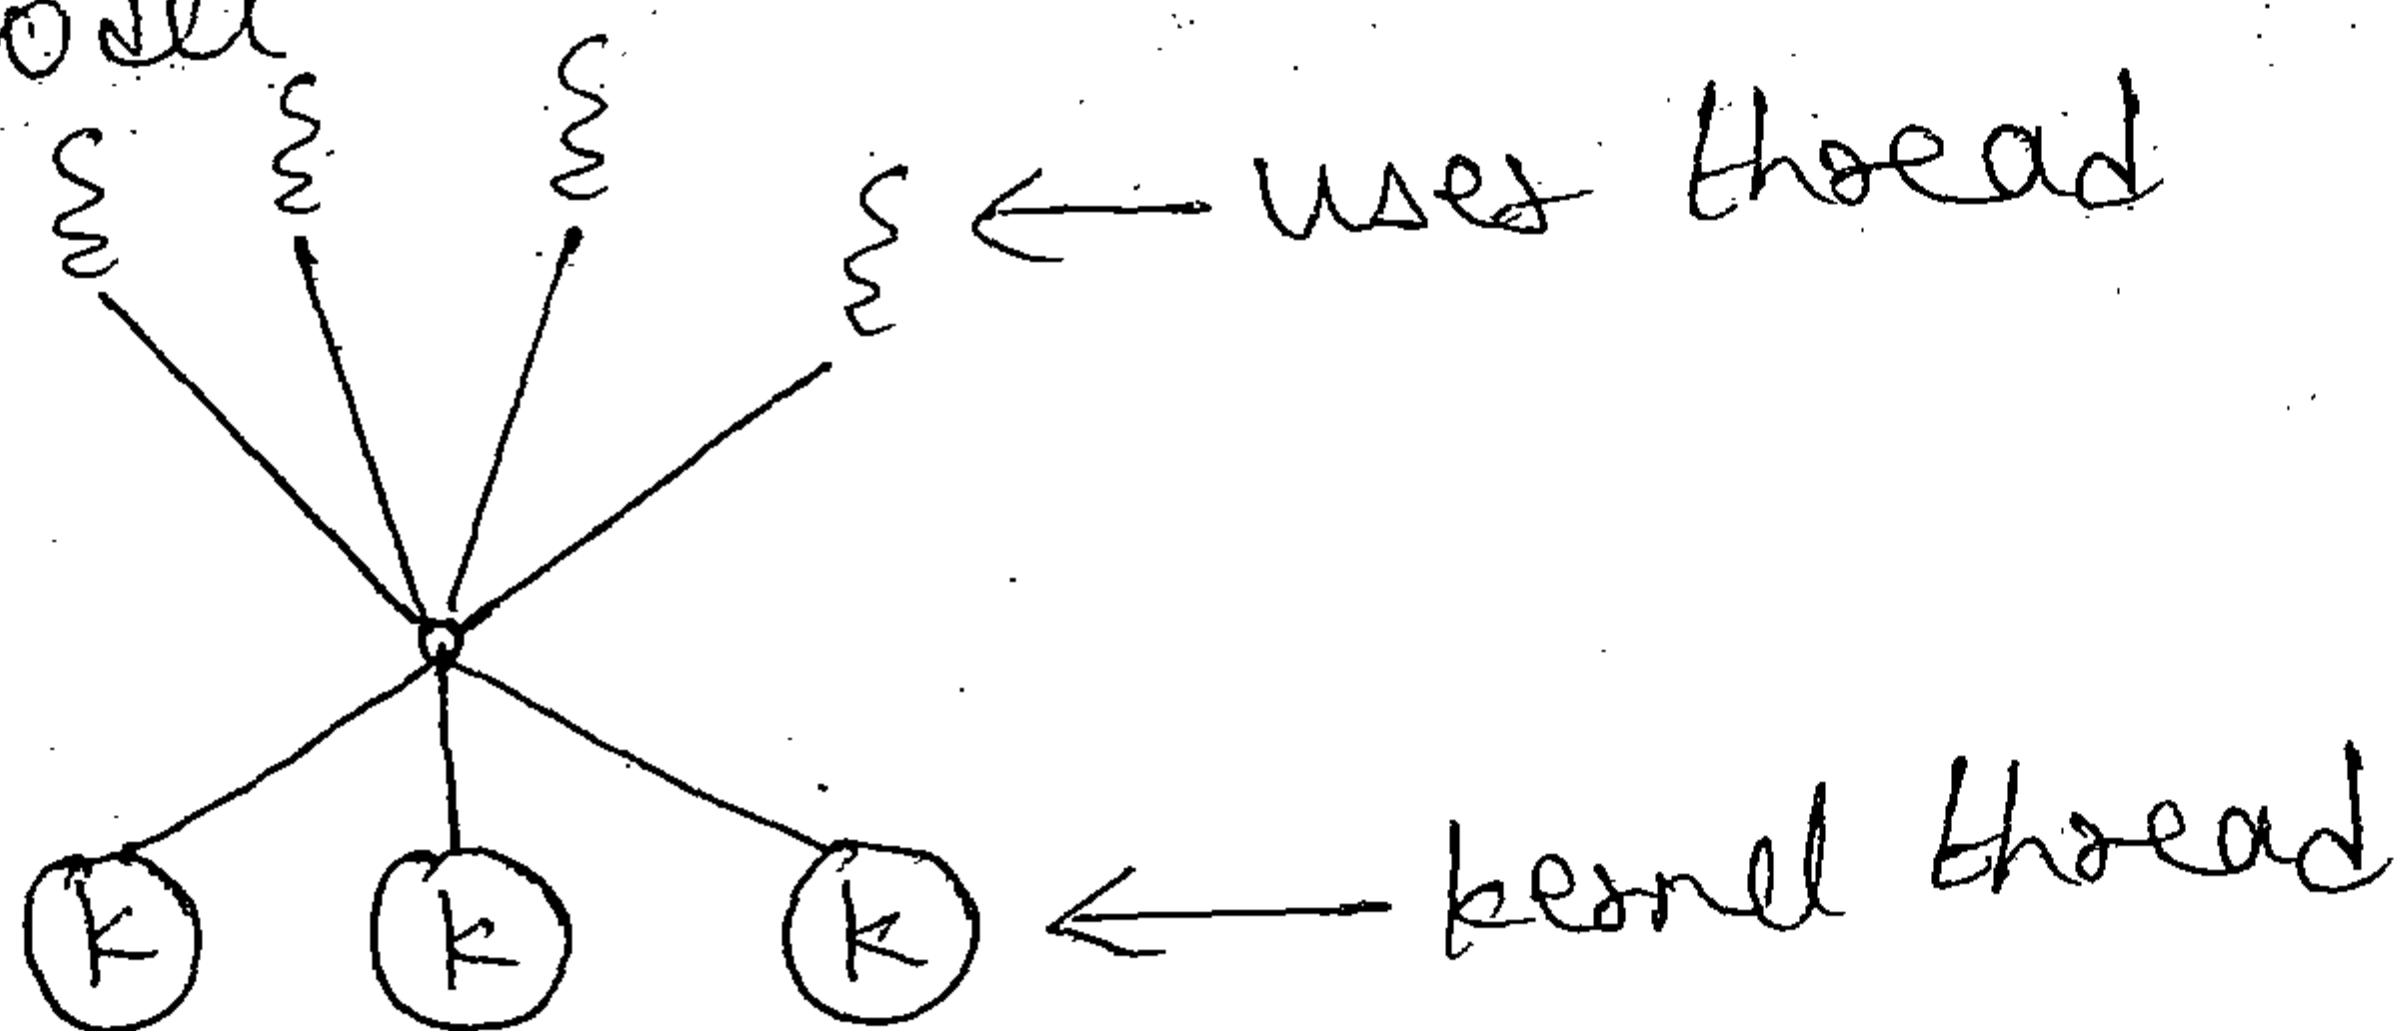
→ Many user-level threads mapped to single kernel thread.
 → Examples: Solaris Green Threads, GNU Portable Threads.

One-to-One Model




- Each user-level thread maps to kernel thread.
- It provides more concurrency than many-to-one model by allowing another thread to run when a thread makes a blocking system call.
- Examples: Windows NT/XP/2000, Linux, Solaris 9 & later.

Many-to-Many Model



- Multiplexes ~~many~~ many user-level threads to a smaller (or) equal number of kernel threads.
- Allows the OS to create a sufficient number of kernel threads.
- One popular variation on the many-to-many model still multiplexes many user-level threads to a smaller (or) equal number of kernel threads but also allows a user-level thread to be bound to a kernel thread — Two-level model



- Solaris prior to version 9.
- Two-level model is supported by OS: IRIX, HP-UX and Tru64 UNIX.
- Windows NT/2000 with the Thread/Fiber package.

① Thread Libraries

- A thread library provides the programmer with an API for creating and managing threads.
- Two primary ways of implementing:
 - * Library entirely in user space.
 - * Kernel-level library supported by the OS.
- Three main thread libraries:
 - 1) POSIX Pthreads
 - 2) Win32
 - 3) Java

Pthreads

- May be provided either as user-level (or) kernel-level.
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization.
- API specifies behavior of the thread library, implementation is up to development of the library.
- Common in UNIX operating systems (Solaris, Linux, Mac OS X).

Win32 Threads

- The technique for creating threads using the Win32 thread library is similar to the Pthreads technique in several ways.

Java Threads

- Threads are the fundamental model of program execution in a Java program.

- Java threads are managed by the JVM.
- Two techniques for creating threads in a Java program.
 - 1) Create a new class that is derived from the Thread class and to override its run() method.
 - 2) Define a class that implements the Runnable interface.

The Runnable interface is defined as follows:

```
public interface Runnable
{
    public abstract void run();
}
```

- Creating a Thread object does not specifically create the new thread.
- start() method that creates the new thread.
- Calling the start() method for the new object does two things:
 - (i) It allocates memory and initializes a new thread in the JVM.
 - (ii) It calls the run() method making the thread eligible to be run by the JVM.

⑦ Threading Issues

The fork() and exec() system calls

- The semantics of the fork() and exec() system calls change in a multithreaded program.
- If one thread in a program calls fork(), does the new process duplicate all threads, (or) is the new process single-threaded?

* Some UNIX systems have chosen to have two versions of fork().

57

→ If a thread invokes the `exec()` system call, the program specified in the parameter to `exec()` will replace the entire process — including all threads.

Cancellation

- Thread cancellation is the task of terminating a thread before it has completed.
 - A thread that is to be canceled is often referred to as: Target Thread.
 - Cancellation of target thread may occur in two different scenarios:
 - 1) Asynchronous cancellation — One thread immediately terminates the target thread.
 - 2) Deferred cancellation — The target thread periodically checks whether it should terminate, allowing it an opportunity to terminate itself in an orderly fashion.
-

Signal Handling

- A signal is used in UNIX systems to notify a process that a particular event has occurred.
- A signal handler is used to process signals. All signals follow:
 - (i) A signal is generated by the occurrence of a particular event.
 - (ii) A generated signal is delivered to a process.
 - (iii) Once delivered, the signal must be handled.
- A signal may be handled by one of two possible handlers:
 - 1) A default signal handler.
 - 2) A user-defined signal handler.

→ Options:

- * Deliver the signal to the thread to which the signal applies.
- * Deliver the signal to every thread in the process.
- * Deliver the signal to certain threads in the process.
- * Assign a specific thread to receive all signals for the process.

→ Windows does not explicitly provide support for signals, they can be emulated using asynchronous procedure calls (APCs).

Thread Pools

- Create a number of threads in a pool where they await work.
- Thread pools offer these benefits:
 - 1) Servicing a request with an existing thread is usually faster than waiting to create a thread.
 - 2) A thread pool limits the number of threads that exist at any one point.

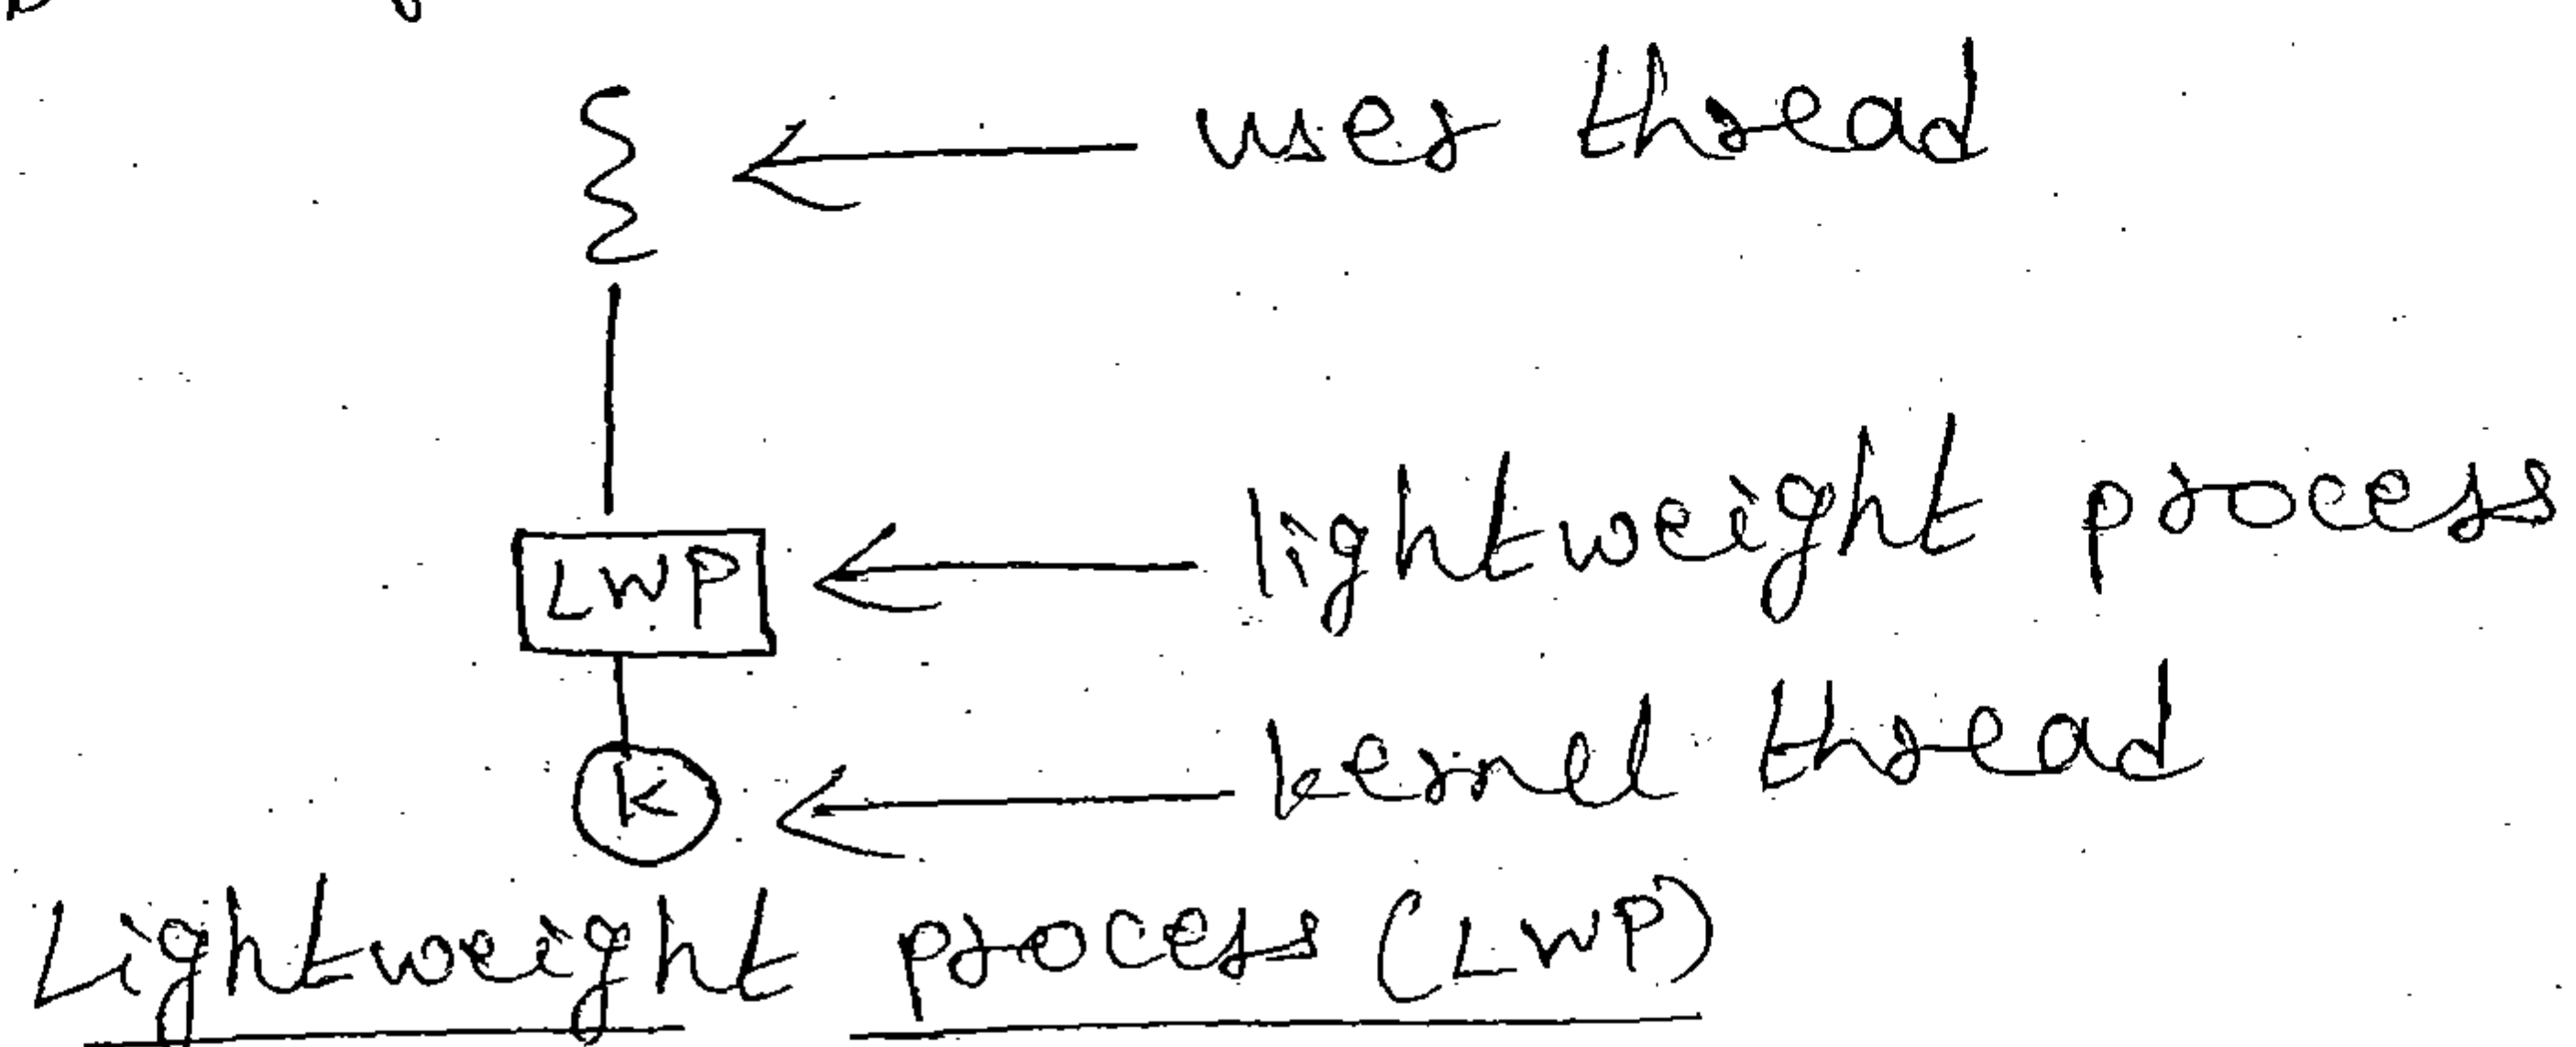
This is particularly important on systems that cannot support a large number of concurrent threads.

= Thread-Specific Data

- Allows each thread to have its own copy of data.
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Most thread libraries - including Win32 and pthreads provide some form of support for thread-specific data.

Scheduler Activations

- Both many-to-many and two-level models require communication to maintain the appropriate number of kernel threads allocated to the application.
- Scheduler activations provide upcalls - a communication mechanism from the kernel to the thread library.
- This communication allows an application to maintain the correct number of kernel threads.



Chapter 5: Process Scheduling

- Objectives:
- * To introduce CPU scheduling, which is the basis for multiprogrammed operating systems.
 - * To describe various CPU-scheduling algorithms.
 - * To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system.

① Basic Concepts.

- In a single-processor system, only one process can run at a time.
- The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
- The problem of determining when processors should be assigned and to which processes is called process scheduling.

- 60
- > When more than one process is runnable, the OS must decide which one first. The part of the OS concerned with this decision is called the scheduler.
- * Algorithm it uses is called the scheduling algorithm.

≡ CPU-I/O Burst Cycle

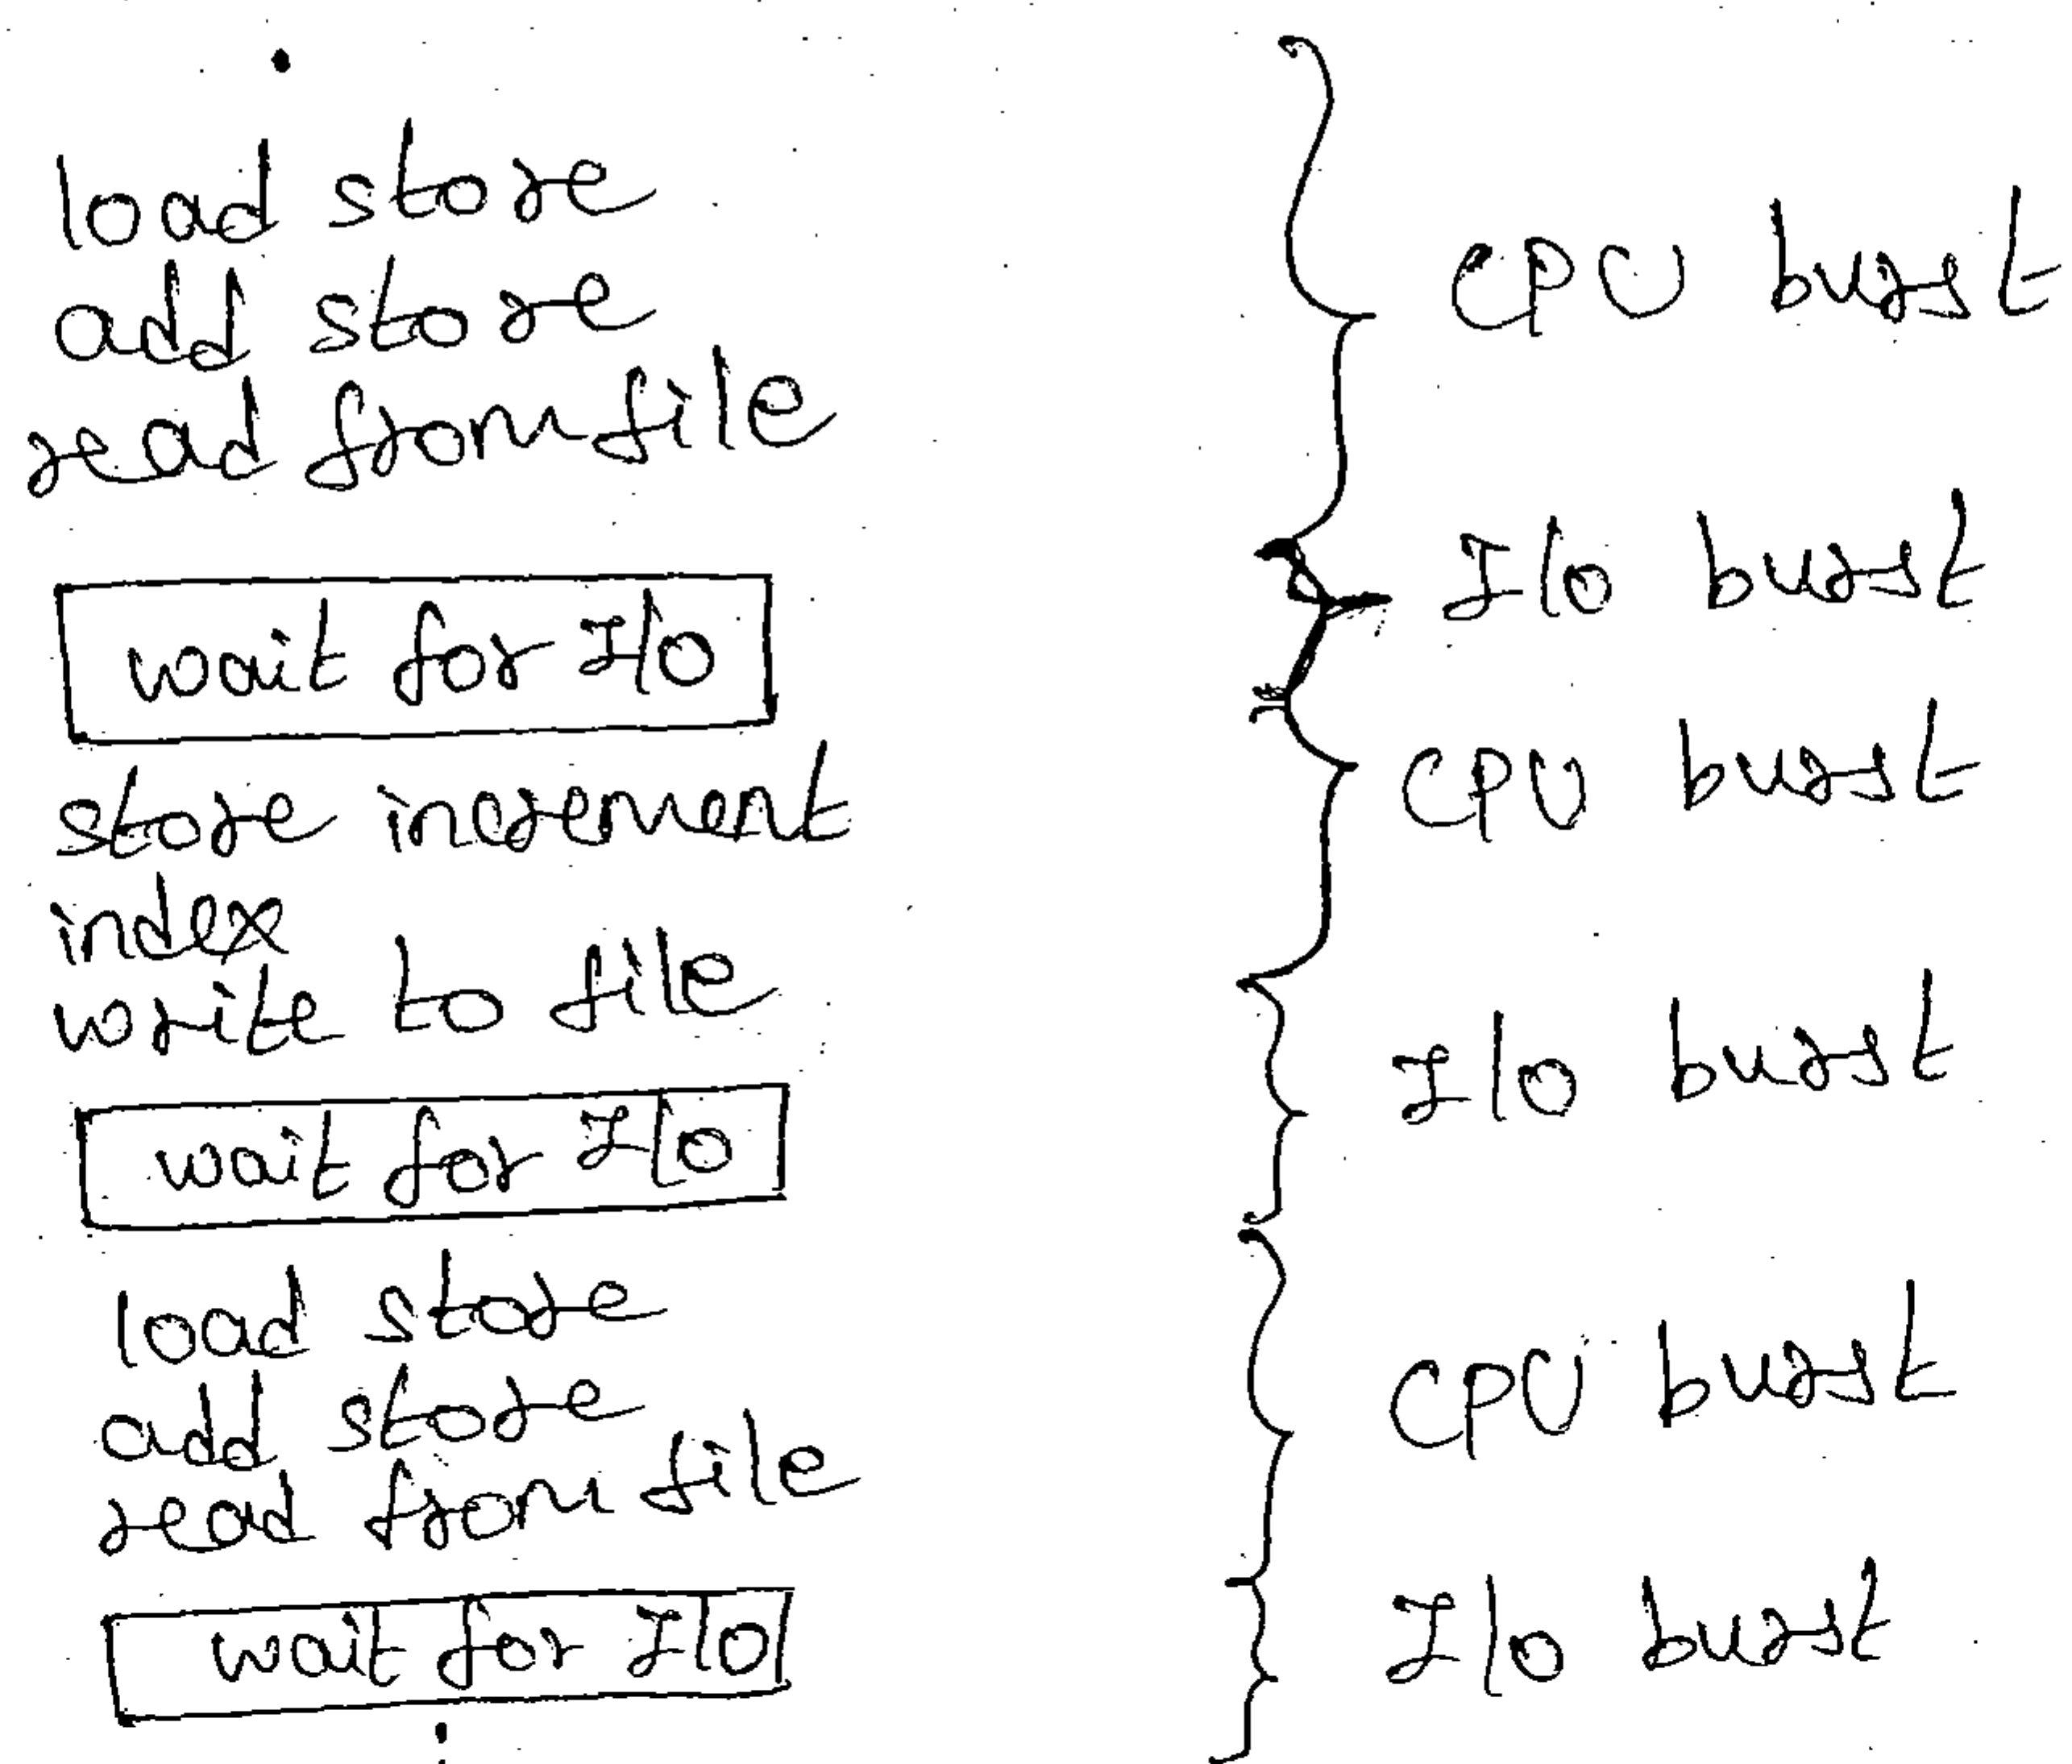
> Process execution consists of a cycle of CPU execution and I/O wait.

> Processes alternate between these two states:

(i) process execution begins with a CPU burst.

(ii) followed by an I/O burst.

Alternating sequence of CPU & I/O bursts



= CPU Scheduler

> selects from among the processes in ready queue, and allocates the CPU to one of them.

* Queue may be ordered in various ways.

> A ready queue can be implemented as a FIFO queue, a priority queue, a tree (or) simply an unordered linked list.

Preemptive Scheduling

→ CPU-scheduling decisions under four circumstances:

1) When a process switches from the running state to the waiting state.

2) When a process switches from the running state to the ready state.

3) When a process switches from the waiting state to the ready state.

4) When a process terminates.

→ Scheduling under 1) and 4) is nonpreemptive (or) cooperative

→ All other scheduling is preemptive.

* Consider access to shared data.

* Consider preemption while in kernel mode.

* Consider interrupts occurring during crucial OS activities.

Dispatcher

→ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler, this involves:

* Switching context

* Switching to user mode

* Jumping to the proper location in the user program to restart that program.

→ Dispatch latency - time it takes for the dispatcher to stop one process and start another running.

⑩ Scheduling Criteria

→ The Criteria include the following:

- (i) CPU utilization - keep the CPU as busy as possible.
- (ii) Throughput - Number of processes that complete their execution per time unit.
- (iii) Turnaround time - Amount of time to execute a particular process.
- (iv) Waiting time - Amount of time a process has been waiting in the ready queue.
- (v) Response time - Amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

⑪ Scheduling Algorithms


→ scheduling algorithm optimization Criteria

- * Max CPU utilization
- * Max throughput
- * Min turnaround time
- * Min waiting time
- * Min response time

First-Come, First served (FCFS) scheduling

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	>

* Suppose that the processes arrive in the order: P_1, P_2, P_3 .
The Gantt chart for the schedule is:



* Gantt chart - A bar chart that illustrates a particular schedule, including the start and finish times of each of the participating processes.


* Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$ milliseconds.

* Average waiting time: $(0 + 24 + 27)/3 = 17$ milliseconds.

* Turnaround time for $P_1 = 24$; $P_2 = 27$; $P_3 = 30$ milliseconds.

* Average turnaround time: $(24 + 27 + 30)/3 = 27$ milliseconds.

* Suppose that the processes arrive in the order: P_2, P_3, P_1 .
The Gantt chart for the schedule is:



* Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$ milliseconds.

* Average waiting time: $(6 + 0 + 3)/3 = 3$ milliseconds.

* Turnaround time for $P_1 = 30$; $P_2 = 3$; $P_3 = 6$ milliseconds.

* Average turnaround time: $(30 + 3 + 6)/3 = 13$ milliseconds.

→ much better than previous case.

→ Convoy effect - short process behind long process.

* Consider one CPU-bound and many I/O-bound processes.

Shortest-Job-First (SJF) scheduling

64

→ Also called as Shortest-Process-Next (SPN).

→ Non-preemptive discipline

→ Associate with each process the length of its next CPU burst.

* Use these lengths to schedule the process with the shortest time.

→ SJF is optimal - gives minimum average waiting time for a given set of processes.

* The difficulty is knowing the length of the next CPU request.

→ Two schemes:

* Non-preemptive - once CPU given to the process it cannot be preempted until completes its CPU burst.

* Preemptive - if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.

→ This scheme is known as Shortest-Remaining-Time-First (SRTF) scheduling.

<u>Process</u>	<u>Burst Time</u>
P ₁	6
P ₂	8
P ₃	7
P ₄	3

→ SJF Scheduling, Gantt Chart:

P ₄	P ₁	P ₃	P ₂
----------------	----------------	----------------	----------------

* Waiting time for $P_1=3$; $P_2=16$; $P_3=9$; $P_4=0$ milliseconds.⁶⁵

* Average waiting time: $(3+16+9+0)/4 = 7$ milliseconds.

* Turnaround time for $P_1=9$; $P_2=24$; $P_3=16$; $P_4=3$ milliseconds.

* Average turnaround time: $(9+24+16+3)/4 = 13$ milliseconds.

→ The next CPU burst is generally predicted as an exponential average of the measured lengths of previous CPU bursts.

→ Exponential average formula:

(i) t_n = length of the n th CPU burst.

(ii) T_{n+1} = predicted value for the next CPU burst.


(iii) α , $0 \leq \alpha \leq 1$ define,

$$T_{n+1} = \alpha t_n + (1-\alpha)T_n.$$

→ Example for shortest-remaining-time-first (SRTF) scheduling.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

* Preemptive SRTF Gantt chart:



* Waiting time for $P_1=(10-1)$; $P_2=(1-1)$; $P_3=(17-2)$; $P_4=(5-3)$ milliseconds.

* Average waiting time: $[(10-1) + (1-1) + (17-2) + (5-3)]/4 = 26/4 = 6.5$ milliseconds.

* Non-preemptive SRTF scheduling would result in an average waiting time of 7.75 milliseconds.

* Turnaround time for $P_1=17$; $P_2=5$; $P_3=26$; $P_4=10$ milliseconds.

* Average turnaround time: $(17+5+26+10)/4 = 58/4 = 14.5$ milliseconds.

Priority Scheduling

- SJF algorithm is a special case of the general priority scheduling algorithm.
- A priority number (integer) is associated with each process.
- CPU is allocated to the process with the highest priority.
(smallest integer = highest priority)
- Can be either Preemptive (or) non-preemptive.
- A preemptive priority scheduling algorithm will preempt the CPU if the priority of the newly arrived process is higher than the priority of the currently running process.
- A nonpreemptive priority scheduling algorithm will simply put the new process at the head of the ready queue.


Major problem with priority scheduling is Indefinite blocking (or) Starvation.

* Low priority processes may never execute.

Solution - Aging: as time progresses increase the priority of the process.

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P ₁	10	3
P ₂	1	1
P ₃	2	4
P ₄	1	5
P ₅	5	2

* Priority scheduling Gantt chart:



* Waiting time for P₁ = 6; P₂ = 0; P₃ = 16; P₄ = 18; P₅ = 1 milliseconds.

* Average waiting time: $(6 + 0 + 16 + 18 + 1) / 5 = 41 / 5 = 8.2$ milliseconds.

* Turnaround time for P₁ = 16; P₂ = 1; P₃ = 18; P₄ = 19; P₅ = 6 milliseconds.

* Average Turnaround time: $(16 + 1 + 18 + 19 + 6) / 5 = 60 / 5 = 12$ milliseconds.

Round-Robin Scheduling

→ Designed especially for time-sharing systems.

→ Similar to FCFS scheduling, but preemption is added to enable the system to switch between processes.

→ Small unit of time called time quantum (q).

→ After this time has elapsed, the process is preempted and added to the end of the ready queue.

→ If there are ' n ' processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once.

* No process waits more than $(n-1)q$ time units.

→ Timer interrupts every quantum to schedule next process.

→ Performance:


* q large \Rightarrow FIFO

* q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high.

→ Example of RR with Time Quantum, $q = 4$.

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

* The Gantt chart is:



* Waiting time for $P_1 = 0 + (10 - 4)$; $P_2 = 4$; $P_3 = 7$ milliseconds.

* Average waiting time: $\frac{[(10 - 4) + 4 + 7]}{3} = \frac{17}{3} = 5.66$ milliseconds.






* Turnaround time for $P_1 = 30$; $P_2 = 7$; $P_3 = 10$ milliseconds.

* Average Turnaround time: $\frac{(30 + 7 + 10)}{3} = \frac{47}{3} = 15.66$ milliseconds.

→ Typically, higher average turnaround than SJF, but better response.

* Multilevel Queue scheduling

highest priority



lowest priority

multilevel queue scheduling

- Ready queue is partitioned into separate queues, example.
 - * Foreground (interactive) processes
 - * Background (batch) processes.
- Process permanently in a given queue.
- Each queue has its own scheduling algorithm:
 - * foreground - RR
 - * Background - FCFS
- Scheduling must be done between the queues:
 - * Fixed priority scheduling - possibility of starvation
 - * Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes i.e., 80% to foreground in RR.
 - * 20% to background in FCFS.

Multilevel Feedback Queue Scheduling

- A process can move between the various queues; aging can be implemented this way.
- multilevel-feedback queue scheduler defined by the following parameters:
 - (i) Number of queues
 - (ii) scheduling algorithms for each queue
 - (iii) method used to determine when to upgrade a process.
 - (iv) method used to determine when to demote a process
 - (v) method used to determine which queue a process will enter when that process needs service.

→ Example:

- * Three queues:
 - Q_0 - RR with time quantum 8 milliseconds.

→ Q_1 - RR time quantum 16 milliseconds.

→ Q_2 - FCFS

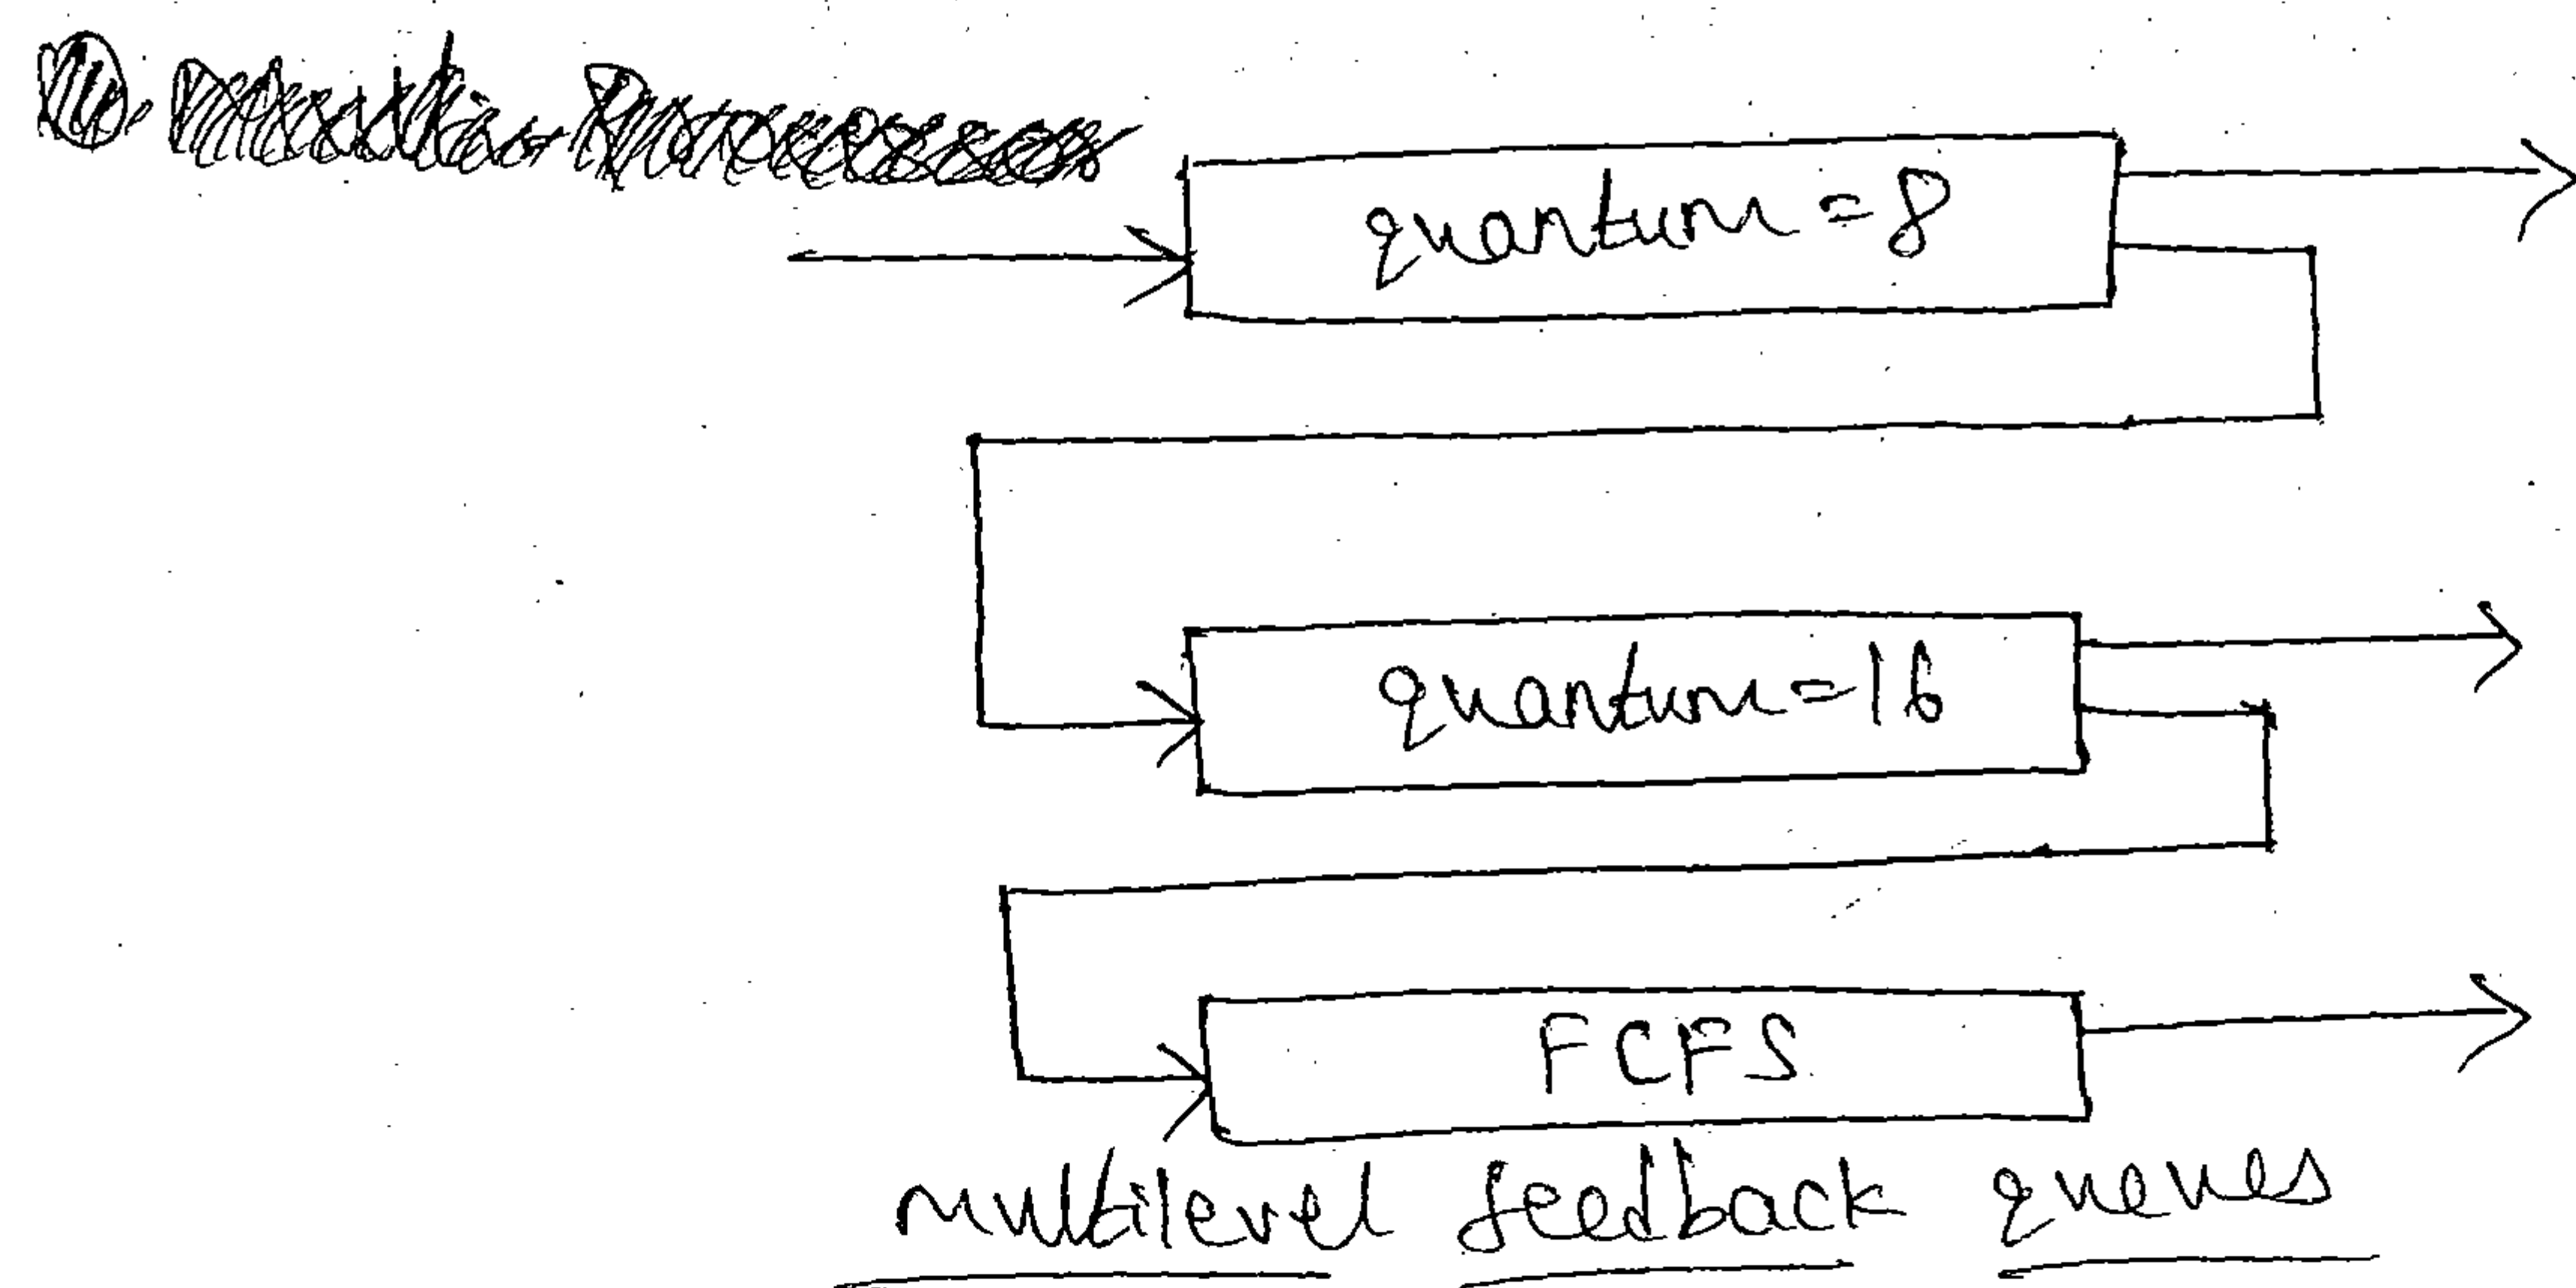
* Scheduling

→ A new job enters queue Q_0 which is served FCFS.

- when it gains CPU, job receives 8 milliseconds.
- If it does not finish in 8 milliseconds, job is moved to queue Q_1 .

→ At Q_1 , job is again served FCFS & receives 16 additional milliseconds.

- If it still does not complete, it is preempted and moved to queue Q_2 .



Multi-processor scheduling

→ CPU scheduling more complex when multiple CPUs are available.

→ Homogeneous processors - within a multiprocessor.

→ Asymmetric multiprocessing - only one processor accesses the system data structures, reducing the need for data sharing.

→ Symmetric multiprocessing (SMP) - each processor is self-scheduling,

all processors can access system data structures.

- private queue of ready processes.

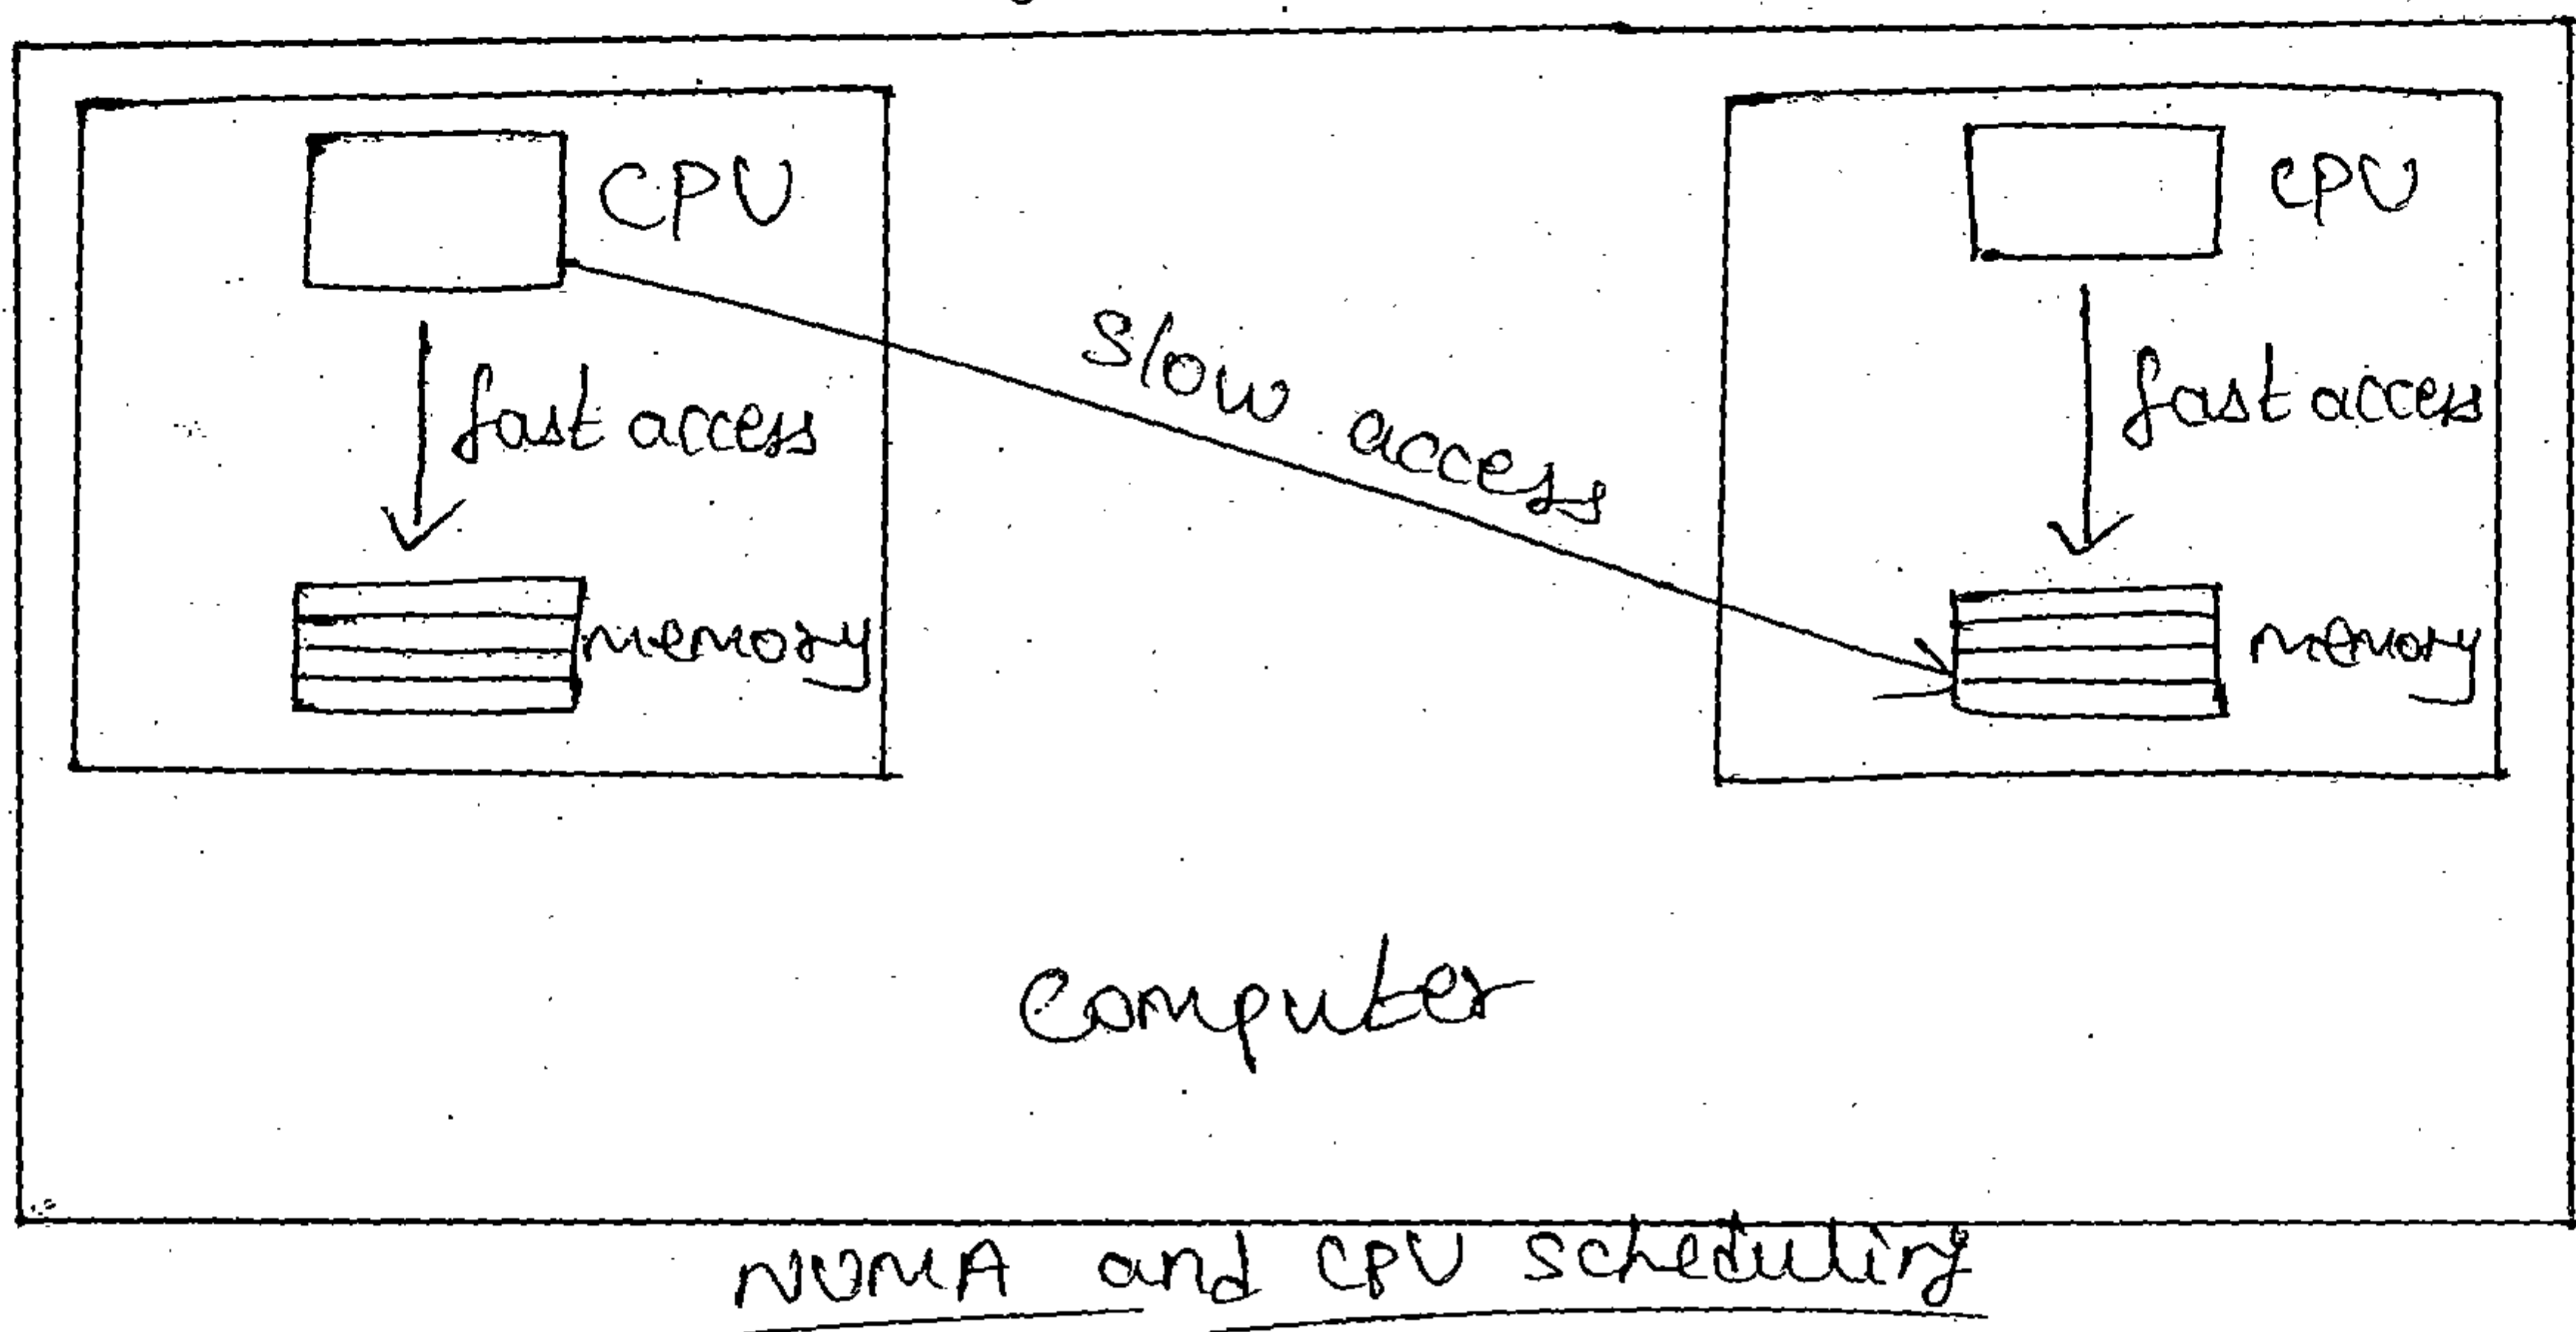
* Currently, most common.

→ Processor affinity - process has affinity for processor which it is currently running.

* Soft Affinity

* Hard Affinity

* Variations including processor sets.



→ Load Balancing - attempts to keep the workload evenly distributed across all processors in an SMP system.

→ Approaches - push migration and pull migration.

→ Multicore Processors.

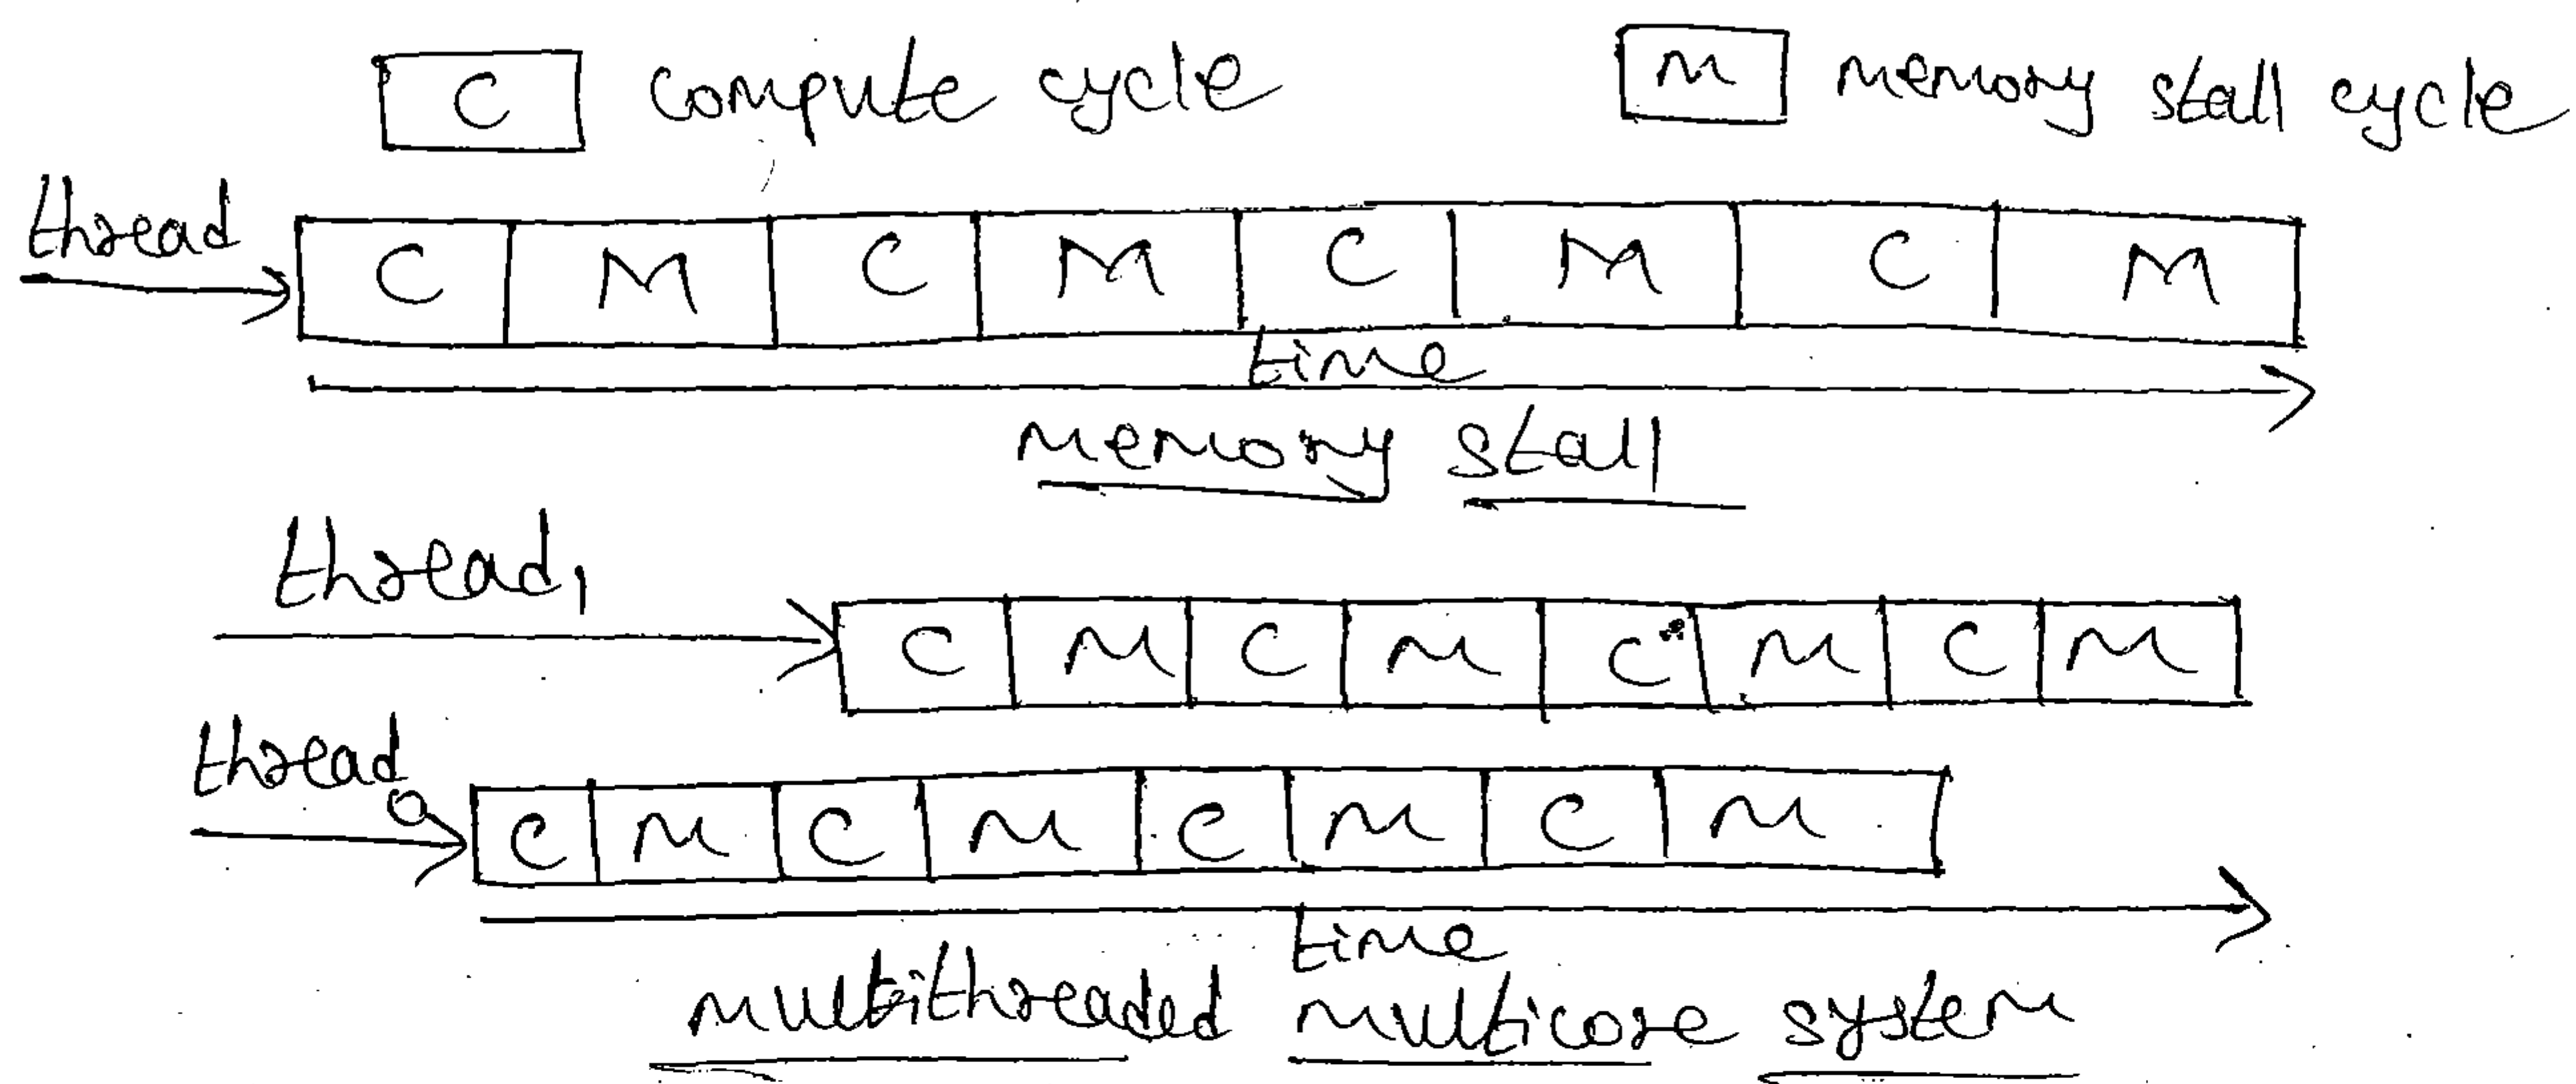
* Recent trend to place multiple processor cores on same physical chip.

* Faster and consumes less power.

* Multiple threads per core also growing

→ Takes advantage of memory stall to make progress on another thread while memory retrieve happens.

* Two ways to multithread a processor: Coarse-grained and



→ Virtualization and Scheduling

- * Virtualization software schedules multiple guests onto CPU(s).
- * Each guest doing its own scheduling
 - Not knowing it doesn't own the CPUs
 - Can result in poor response time.
 - Can effect time-of-day clocks in guests.
- * Can undo good scheduling algorithm efforts of guests.

⑦ Thread Scheduling

- Distinguishing between user-level and kernel-level threads.
- On OS that support them, it is kernel-level threads.
- Not processes - that are being scheduled by the OS.
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on lightweight process (LWP).
- * known as process-contention scope (PCS) since scheduling competition is within the process.
- * Typically done via priority set by programmer.
- kernel thread scheduled onto available CPU is system-contention-scope (SCS) - competition among all threads in system.

→ Pthread Scheduling

* API allows specifying either PCS (or) SCS during thread creation.

⇒ PTHREAD_SCOPE_PROCESS schedules threads using PCS scheduling.

⇒ PTHREAD_SCOPE_SYSTEM schedules threads using SCS scheduling.

* The pthread_t PC provides two functions for getting and setting the contention scope policy:

⇒ pthread_attr_setscope(pthread_attr_t *attr, int scope)

⇒ pthread_attr_getscope(pthread_attr_t *attr, int *scope).

The first part of the document discusses the importance of maintaining accurate records of all transactions. It emphasizes that every entry, no matter how small, should be carefully documented to ensure the integrity of the financial data. This includes recording dates, amounts, and the nature of the transactions.

The second part of the document outlines the procedures for reconciling the accounts. It states that a thorough reconciliation should be performed at the end of each month to identify any discrepancies between the recorded transactions and the actual bank statements. Any differences should be investigated and corrected promptly.

The third part of the document provides a detailed breakdown of the expenses incurred during the period. It lists various categories such as salaries, rent, utilities, and supplies, and provides a clear summary of the total costs for each category. This information is crucial for understanding the overall financial performance and for budgeting purposes.

The final part of the document concludes with a statement of the net income or loss for the period. It summarizes the total revenue, total expenses, and the resulting profit or loss. This section is essential for providing a clear picture of the organization's financial health and for making informed decisions about future operations.